

SAP HACKFEST 2026 · THEME 1

SANJEEVANI. The Unified Dossier.

One continuous narrative in four parts: the case, the build, the run and the road beyond. The Master Strategy Dossier and the Global Deep-Dive Research are merged here into a single flow — deduplicated, renumbered and cross-referenced — covering the self-healing pharma cold chain end to end.

Prepared for: **Hackfest Team**
South & East Regional Round · Theme 1

SEPTEMBER 2026 · UNIFIED EDITION

Table of Contents

PART A - The Case: Evidence and Selection	1
1. Executive Verdict: Build SANJEEVANI	2
2. How Supply Chains Break: The Global Evidence	4
2.1 Research method and source discipline	4
2.2 Six disruptions that define the problem	4
2.3 The pattern: sensing is solved, deciding is not	6
2.4 The pharmaceutical leak, quantified	6
2.5 Where excursions actually happen	7
2.6 Climate, complexity and the money at stake	7
3. The Self-Healing Thesis: Monitoring Is Saturated, Recovery Is Empty	8
4. Market and Competitor Landscape	9
4.1 Capability-gap matrix	10
5. Community and Social Signals: What Practitioners Are Actually Saying	12
6. The Research Frontier: Solved, Open and Claimable	13
6.1 Solved versus open	15
6.2 The five claimable gaps	16
7. Hugging Face and the Open-Source Build Inventory	17
7.1 Agent orchestration	19
7.2 Simulation, optimization and forecasting assets	19

8. The Regulatory Stack: Compliance as Architecture	21
9. India Ground Reality: Exports, the 1% Problem and eVIN	22
9.1 Buyer pain signals and the India angle	24
10. Patent and IP Landscape	24
10.1 Prior art found	24
10.2 White-space hypotheses (requiring professional patentability review)	25
11. SAP Ecosystem Fit: Riding the Vendor's Own Roadmap	26
11.1 SAP has already written our pitch frame	26
12. Fifteen Concepts and the Final Selection	28
12.1 Closed-loop execution concepts	28
12.2 Network intelligence concepts	29
12.3 Platform and infrastructure concepts	31
12.4 Why the winner beats #2 and #3	35
PART B - The Build: Architecture and Design	36
13. Winning Project Deep Dive: Personas and Problem Statements	37
13.1 Problem statement (rubric-ready) Problem Statement - one sentence	37
13.2 Primary persona and empathy map	37
13.3 User needs (evidence-backed)	38
13.4 Current workflow vs SANJEEVANI workflow	38
13.5 The three-layer design (MVP / Research / Startup)	39
14. System Architecture: The Closed Loop	40
14.1 Component-to-SAP mapping and demo feasibility	41
14.2 Safety, governance, and "what happens when the AI is wrong"	42

14.3 Sanjivani v2: the lane-level cold-chain view	43
14.4 Escalation matrix: autonomy with boundaries	44
15. Agent Specifications and the Full Roster	45
15.1 FMEA snapshot: failure modes and mitigations	47
16. Data Architecture	50
16.1 Core entities (knowledge graph schema sketch)	50
16.2 Data flows and storage decisions	50
17. Technology Stack, Licensing and Competitive Positioning	51
17.1 Where each player stands on the loop	51
17.2 The SAP stack on student and trial licenses	52
17.3 The SAP student stack reality check	53
17.4 Final stack with justification	53
17.5 Competitive advantage (Part Q, answered directly)	54
17.6 IF I WERE YOU - the brutally honest closing	55
PART C - The Run: Winning the Two Rounds	56
18. Round 1: The Ten-Minute Pitch	57
19. Round 2: The Live Demo - The Ten-Minute Recovery	58
19.1 Quantitative comparisons to show (and defend)	60
19.2 Fallback demo plan	60
20. The Q&A; War Room: Thirty Questions, Strong Answers	61
21. Evaluation Framework and Risk Register	64
21.1 Four-layer metric framework	64

21.2 Project risk register (beyond agent FMEA)	65
22. Implementation Plan and Team Assignments	66
22.1 Round 1 submission plan (by September 22)	66
22.2 Prototype sprint (September 23-29) and finale (September 30)	67
22.3 Team assignments (4-5 members)	68
23. Round 1 Submission Templates (Fill-Ready)	68
23.1 Empathy map - Meera Iyer (ready to paste)	68
23.2 System failure-mode map (current state)	69
23.3 Agent-human interaction map	69
23.4 Business Model Canvas - one-page version	70
PART D - Beyond the Hackfest: Research, Patent, Startup	71
24. Research Roadmap: RecoveryBench-12	72
24.1 Research question and hypothesis	72
24.2 RecoveryBench-12: the benchmark contribution	72
24.3 Metrics, ablations and paper plan	73
25. Patent and IP Roadmap	73
25.1 Student-executable filing strategy	74
26. Startup Plan and Business Model	75
26.1 Company concept and ICP	75
26.2 Business Model Canvas	75
26.3 Go-to-market and the 1/3/5-year roadmap	76
26.4 "If the Hackfest ended tomorrow, what could we sell next week?"	76
26.5 Moat ranking	76

27. Financial Model: Conservative, Base, Aggressive 77

27.1 Customer ROI equation (the value model that prices the product) 78

28. Evidence Library and Glossary 79

PART A

PART A - The Case: Evidence and Selection

Why this project, proven with evidence. The chapters in Part A establish the problem, the competitive whitespace, the research and regulatory grounding, the India rationale, and the selection decision that produced SANJEEVANI from a field of fifteen candidates. Content from both source volumes is merged and deduplicated: the global deep-dive evidence carries the argument, and the master strategy supplies the selection machinery.

Chapter 1 Executive Verdict: Build SANJEEVANI

Chapter 2 How Supply Chains Break: The Global Evidence

Chapter 3 The Self-Healing Thesis: Monitoring Is Saturated, Recovery Is Empty

Chapter 4 Market and Competitor Landscape

Chapter 5 Community and Social Signals

Chapter 6 The Research Frontier: Solved, Open and Claimable

Chapter 7 Hugging Face and the Open-Source Build Inventory

Chapter 8 The Regulatory Stack: Compliance as Architecture

Chapter 9 India Ground Reality: Exports, the 1% Problem and eVIN

Chapter 10 Patent and IP Landscape

Chapter 11 SAP Ecosystem Fit: Riding the Vendor's Own Roadmap

Chapter 12 Fifteen Concepts and the Final Selection

1. Executive Verdict: Build SANJEEVANI

This unified dossier merges two earlier volumes into one continuous narrative: the **Master Strategy Dossier**, which ran the full selection process from ninety-one structured searches to a fifteen-concept portfolio and a weighted decision, and the **Global Deep-Dive Research**, which stress-tested the winning concept against worldwide evidence across players, communities, Hugging Face models, papers, regulation and India’s ground reality. Overlapping sections were deduplicated, the strongest material from each volume was kept, and everything is renumbered into a single flow. The verdict itself can be stated before any evidence: among pharma cold chains, disaster relief logistics, tariff-exposed networks and agri-food supply chains, the pharmaceutical cold chain earned the highest weighted score, and the team locked it in - not as another monitoring dashboard, but as **Project SANJEEVANI**, a self-healing supply chain in which software agents sense a disruption early, convert it into simulated scenarios, generate an optimized recovery plan, and execute it through SAP systems after a human approval gate.



Project name	SANJEEVANI - The Self-Healing Pharma Supply Chain Network
One-line description	A governed multi-agent system that senses supply chain disruptions early, converts them into simulated scenarios, generates an optimized recovery plan, and executes it through SAP systems after a human approval gate - purpose-built for life-critical pharmaceutical supply chains.
Why this problem matters	Disruption is no longer an anomaly; it is the operating environment. 80% of manufacturers experienced a significant disruption within 12-18 months, and the profit hit reached 13%. In pharma, a delayed shipment is not a lost sale - it is a missed dose of a life-saving medicine. India ships \$27.9bn of pharmaceuticals annually with 70-90% API dependency on a single neighbor, making its supply chains structurally exposed.
Why now	Three forces converge in 2026: (1) structural disruption (Red Sea re-routing, tariff volatility, canal droughts) has become permanent; (2) SAP launched Autonomous Supply Chain Management at Sapphire 2026 (6 Joule Assistants, 60+ AI agents) - the vendor itself says the operating model is "Sense, Decide, Act faster"; (3) Gartner projects agentic SCM spend growing 25x by 2030 while warning that over 40% of agentic projects will be cancelled by 2027 - governance, not more agents, is the differentiator.

Why Agentic AI	Disruptions propagate faster than human planning cycles. A single alert triggers dozens of interconnected decisions (re-route, re-source, re-allocate, re-file compliance) that currently take days-to-weeks of manual coordination. Only a team of specialized cooperating agents - sensing, scenario, impact, logistics, compliance, audit - can compress that to minutes while staying inside policy. This maps 1:1 to the official theme agents: Sensing & Scenario, Logistics & Inventory, Compliance & Human-in-the-Loop.
Why SAP	The system is designed against released SAP surfaces: IBP scenario and key-figure OData, TM freight order/booking APIs, Ariba Supplier Risk / Risk Exposure APIs, HANA Cloud vector + knowledge graph engines, Joule/Generative AI Hub for LLM calls. SAP partners (Everstream, Prewave, Resilinc) already sell "risk feeds into IBP" - our architecture is the SAP-certified pattern extended one step further into closed-loop execution.
Why we can win	Round 1 judges score Persona, Problem, User Needs, Solution Alignment. Our persona (a real, evidenced planner persona at an Indian pharma exporter), our problem (quantified alert-to-recovery gap), and our solution (a visible closed loop) align to the rubric with unusual precision. For the finale, the demo is a live 10-minute disruption-to-recovery loop - memorable, explainable and SAP-native.
Why customers would pay	Every customer conversation anchors on avoided loss: stockouts, expedited freight, excursion write-offs, demurrage, and planner overtime. A mid-size pharma exporter losing one biologics shipment (up to 20% of biologics shipments fail on temperature) can justify a \$15k-50k ACV from a single prevented incident.
Why competitors cannot easily copy it	Sensing is crowded (Everstream, Prewave, Resilinc, Interos); visibility is commoditizing (project44, FourKites). Nobody owns the post-alert layer: simulate, decide, get approval, execute with audit. Our moat compounds through a proprietary decision-outcome dataset (which recovery plans actually worked), solver-guardrail architecture (patent hypothesis), and GDP-grade audit trails that competitors selling alerts do not generate.
Research opportunity	A verifiable literature gap: no published system connects weak-signal NLP sensing to LLM-parameterized disruption scenarios, deterministic simulation, solver-based re-planning, resilience scoring and HITL governance in one evaluated pipeline. We can contribute a benchmark (RecoveryBench-12) and an ablation study; see Chapter 17.
IP / patent hypothesis	Three white-space directions surfaced by prior-art screening: (1) guardrailed solver-calling agent architecture with policy-constrained action validation; (2) weak-signal-to-parameterized-scenario pipeline with confidence scoring; (3) GDP-compliance-aware autonomous cold-chain response state machine. These are hypotheses requiring professional patentability review - not claims of patentability.
Startup opportunity	Beachhead: Indian pharma exporters (evidence-backed pain, SAP-ready segment, digital customs rails). Pricing: \$15k-50k mid-market ACV band with success-fee hybrid. Tailwind: \$5-23bn SCRM market growing at 9-14% CAGR, plus Gartner's agentic curve. The hackfest prototype is deliberately architected as Layer 1 of the product (Chapter 10).

Table 1 - Executive verdict summary (Part A format from the master brief).

The one-sentence strategy SAP has published the destination - "autonomous, self-healing supply chains" - and the world's vendors sell the alarm clock; SANJEEVANI demonstrates and eventually sells the healing, inside SAP, with a human holding the final pen.

2. How Supply Chains Break: The Global Evidence

RESEARCH METHOD AND SIX CASE STUDIES WITH NUMBERS

2.1 Research method and source discipline

This dossier is built on three parallel research waves executed in September 2026: 28 searches on the SAP ecosystem, 26 on academic and open-source landscapes, and 37 on markets, competitors, patents and disruption case studies - 91 structured searches in total, with raw results archived as JSON evidence files.

Every major claim in this report carries a source and a date, and conclusions are labeled as FACT (directly reported by a cited source), INFERENCE (our analytical conclusion, with reasoning), or HYPOTHESIS (a claim explicitly requiring further validation, e.g., patentability). Where numbers could not be verified across at least one credible source, they are flagged rather than fabricated. Social-media signals are used only as directional color and are never weighted as evidence.

The selection question we asked the evidence is narrow and practical: which supply-chain scenario gives a 4-5 person student team (a) maximum rubric alignment for the September round, (b) a demoable agentic-AI story that no judge has seen from students before, and (c) a defensible wedge that survives contact with the market after the hackfest ends? The case studies below are the demand-side argument; the landscape chapters that follow are the supply-side argument.

2.2 Six disruptions that define the problem




Figure 1 - Documented disruption cost markers (log scale). Sources: J.S. Held via NetSuite (Apr 2026); AlixPartners via CNBC (2021); Sensos/industry estimates (2025); FAO/OpenKnowledge reporting of Drewry benchmark (Jan 2024); ITF-OECD (Mar 2024); DocShipper (Jan 2026).

The Red Sea crisis is the defining case of the decade. From late 2023, attacks forced median container ships onto the Cape of Good Hope, adding 10-14 days of transit; Red Sea transits fell by up to 90%, and the benchmark container freight rate hit \$5,481 per FEU in the week of 23 January 2024 - a 417% spike against baseline. ITF-OECD put the pure re-routing cost at roughly \$300 per 40ft container before insurance premiums, and UNCTAD documented the resulting effective capacity loss through 2024. What failed was not visibility - the world watched the attacks in real time. What failed was decision speed: enterprises re-planned manually across weeks, absorbing inventory bloat and service failures that better scenario automation could have compressed.

Tariff shock is the second structural case. The US average effective tariff rate reached 16.8% by November 2025 (z2data), and the de minimis exemption - duty-free entry for packages under \$800 - ended for all countries on 29 August 2025, forcing every e-commerce and DTC flow through formal customs. McKinsey reported that tariff changes now affect 30% of global supply chain activities, and the US National Foreign Trade Council found 50% of surveyed firms had already rerouted shipments or changed suppliers in response. KPMG's May 2025 C-suite survey (300 executives) recorded 77% adapting. Columbus Consulting's verdict - tariff volatility will persist for years, not months - is precisely the operating assumption a resilience product should encode.

The remaining cases complete the pattern. The Panama Canal drought cut transits to historic lows before recovering only by mid-2025. The October 2024 US East and Gulf Coast port strike lasted three days but produced a backlog that Xeneta said would take weeks to clear - because recovery sequencing, not awareness, is the bottleneck. The 2021 semiconductor shortage cost the auto industry \$110-210bn in one year and cut global auto production by 26% in nine months. And COVID-era pharma dependency remains live: India sources roughly 70% of APIs (and 70-90% of key starting materials) from China, so a single regional shutdown propagates directly into the world's medicine cabinet.




Figure 2 - The alert-to-recovery gap: public signals existed days in advance, yet manual recovery took weeks-to-months. Sources: Xeneta (Oct 2024), project44 (Sep 2025), Interos/TechMonitor (Nov-Dec 2024), UNCTAD (2024). Bars are indicative reconstructions from reported facts.

2.3 The pattern: sensing is solved, deciding is not

Across all six cases the anatomy is identical. First, a sudden capacity, route, or regulatory shock. Second, weeks-to-months of manual re-planning (Cape re-routing, port backlog clearance, canal recovery, tariff-induced re-sourcing). Third, cost blowouts and human consequences - missed medicines, higher food prices, wiped-out farmer incomes, the exact harms the official theme text names. Fourth, and most important for this hackfest: the information needed to react well was usually available early. JPMorgan and project44 both published Red Sea analyses in early February 2024, days before most corporate responses matured. The failure was never the absence of a signal; it was the absence of an organized decision machine between the signal and the SAP system of record.

Disruptions propagate faster than humans can coordinate. The delta between alert-time and recovery-time is where supply chains bleed - and where an agentic system earns its keep.

Implication for project selection Do not build another sensing dashboard. The winning wedge is the closed loop after the alert: verify, simulate, optimize, approve, execute, learn. That wedge is also, per Chapter 3, the one no funded player owns - and per Chapter 4, the one no academic group has evaluated end-to-end.

2.4 The pharmaceutical leak, quantified

Start with the scale of the leak, because it is larger than most judges expect. The World Health Organization's long-cited estimate holds that up to half of all vaccines are wasted globally every year, with temperature control and logistics failures a leading cause; industry reporting as recently as November 2025 still repeats the figure, and a parallel 2025 analysis puts annual vaccine wastage from temperature failures alone at roughly one in four doses. Beyond immunization, about twenty percent of all temperature-sensitive pharmaceuticals are damaged by or at least exposed to significant temperature excursions during transport,

and roughly twelve percent of pharmaceutical shipments still register excursions even in an era of ubiquitous data loggers. Sensitivity is brutal at the margin: sources across the 2025-26 trade press note that a one-to-two degree drift can degrade biologics, vaccines and insulin, and that freezing is as lethal as heat - and harder to detect, because a frozen-then-thawed vial looks normal.




Figure 3 - The leak in the chain: product lost or exposed, across four recent global estimates (WHO via UNEP; 2025 industry reports; EROAD; pharmaceutical trade press, 2026).

2.5 Where excursions actually happen

The newest and most strategically useful evidence is about **where** excursions actually happen. A 2026 practitioner study by Kanwal and a wave of 2026 trade coverage from air-cargo specialists converge on the same conclusion: in pharmaceutical air freight, excursions occur most frequently during ground operations - airport dwell, tarmac exposure, customs holds and the handoffs between truck and plane - not in flight. Cooperative logistics guidance explicitly warns against assuming temperature-controlled containers eliminate ground risk and against underestimating inspection dwell times. For a hackathon team this is gold: the failure window we simulate is narrow, nameable, frequent, and - judging by the trade press - still managed by phone calls and spreadsheets.

2.6 Climate, complexity and the money at stake

Two forces make the leak worse, not better. First, climate: a 2025 peer-reviewed review catalogs rising temperatures, floods, cyclones and power outages as direct threats to vaccine potency and cold-chain logistics, and India's repeated heatwaves put the problem on the front page every summer. Second, complexity: a 2026 study monitoring vaccine distribution routes in Ukraine found weaknesses even in corridors operated by experienced agencies, suggesting that no lane is too mature to fail. Meanwhile the money at stake keeps growing - healthcare cold chain logistics is projected to expand from 62.5 billion dollars in 2025 to 95.1 billion by 2030, cold chain monitoring from roughly 35 billion dollars in 2024 toward

120 billion by 2030 on Grand View Research’s broader definition, and about eighteen percent of all biopharma logistics spend already goes to cold chain management. A leak this expensive, in a market this large, is exactly the kind of problem enterprise buyers pay agents to close.

3. The Self-Healing Thesis: Monitoring Is Saturated, Recovery Is Empty

The thesis needs restating precisely, because the words are now fashionable. A **monitoring** system senses and alerts: it tells a human that a pallet is drifting toward its temperature limit. A **self-healing** system closes the loop: it predicts the breach, simulates recovery options against cost, risk and compliance constraints, executes the selected play through carrier and warehouse systems, writes the regulatory audit trail, and learns from every event. The six verbs - sense, predict, simulate, act, comply, learn - are the entire architecture. Monitoring is saturated because sensors are cheap and every vendor ships them; recovery is empty because acting autonomously requires the harder combination of simulation, integration and auditability that dashboards never needed.

The language of self-healing has nevertheless gone mainstream, which is why the timing evidence matters. IBM’s consulting arm published its playbook for agentic supply chain operating models in April 2025; the SAP-connected sensing firm Sensos argued in June 2025 that AI agents fused with IoT sensing produce self-healing networks; boutique consultancies and Locus.sh spent 2025 and 2026 explaining that agentic AI moves logistics from decision support to execution, with Locus declaring 2026 the year agentic AI proves real return on investment. Most tellingly, SAP’s own July 2026 article quotes its supply chain thought leader foreseeing that "the future of supply chain will be self-healing, automated and driven by AI," with humans still playing a critical role - which is, almost word for word, the human-in-the-loop architecture this project has pitched since the strategy volume.




Figure 4 - Convergence timeline, late 2023 to mid 2026: capital, repositioning, open-source milestones and SAP announcements all point toward autonomous recovery. Compiled from company press releases and research repositories.

Reading the timeline as evidence, four streams converge. Capital flowed into cold-chain intelligence (Controlant's 80 million dollar round in November 2023; PAXAFE's Series A in May 2024). Incumbents repositioned toward decisions rather than visibility (Roambee becoming Decklar in September 2025 explicitly to claim "Real-Time Decision AI"). SAP moved twice in eight months - three Joule supply chain agents in October 2025, then the Autonomous SCM portfolio at Sapphire 2026. And the open-source engine reached production quality: Chronos-Bolt landed on Hugging Face in November 2024, the GIFT-Eval benchmark gave zero-shot forecasting a leaderboard, and NVIDIA's TiRex topped it in May 2025. Yet the decisive fact survives all this motion: a 2026 SSRN review observes that public examples of pharma digital twins remain limited; our Devpost and GitHub scans found no completed self-healing cold-chain demo; and practitioner threads on Reddit still describe alert fatigue and manual corrective calls. The loop remains unclosed in public - and that is precisely the gap a student team can close first.

4. Market and Competitor Landscape

The competitive field sorts cleanly into three tiers plus an adjacent ring. The first tier is **hardware and sensing**: Sensitech (owned by Carrier) fields data loggers, real-time trackers and site monitoring across food and life sciences; Tive sells the Solo line of cellular, WiFi and GPS trackers - including flight-compliant Solo 5G devices and the budget Solo Lite launched in June 2024 - and its own guidance describes excursion response as a human workflow: root cause, then a team that "reroutes the shipment to a qualified storage facility" and calls the carrier. That sentence is the best public evidence that the act stage of the loop is still manual. The second tier is **visibility platforms**: Controlant of Iceland, which raised 80 million dollars in November 2023 and acquired business-intelligence firm Intenta, positions itself as the real-time visibility cloud for high-value, high-risk pharma shipments.

The third tier is the one to watch: **decision-AI challengers** moving deliberately toward our territory. PAXAFE, a Boston-area startup, calls itself a GxP cold chain decision intelligence and quality compliance platform, raised a 9 million dollar Series A in May 2024, argues in its own blog that the industry needs "orchestration, not visibility dashboards," and in July 2026 announced a partnership with packaging giant Peli BioThermal - hardware and AI starting to bundle. Decklar (formerly Roambee) rebranded on September 10, 2025 as "the world's only Real-Time Decision AI platform for supply chains," a direct claim on the same language we use. TraceLink, the pharma serialization network, is scaling "governed AI agents" across multi-enterprise workflows after a transformative 2025. The adjacent ring matters for demo realism: Envirotainer's RAP e2 and Releye containers dominate active air freight and are approved by major airlines; DHL has been consolidating pharma specialty logistics through the CryoPDP and SDS Rx acquisitions; and Zipline's drones deliver roughly 150 essential medicine products across African health systems - a reminder that recovery execution can take physical forms we can simulate.

Competitive scan: the state of the self-healing loop, September 2026

Player (tier)	What they publicly do	What they do not do - the gap we occupy
Sensitech / Carrier (sensing)	Loggers, real-time trackers, site monitoring for food and life sciences	No autonomous recovery; corrective actions remain human workflows
Tive (sensing)	Solo 5G / Solo Lite trackers; lane risk guidance; excursion analytics content	Explicitly documents human rerouting and carrier calls as the corrective step
Controlant (visibility)	Real-time visibility cloud for high-value pharma; \$80M round (2023); BI acquisition	Visibility, alerts and data - not simulation, negotiation or auto-execution
Envirotainer (containers)	Active RAP e2 / Releye containers; airline approvals; RAP network	Container-level protection only; no lane-level decision autonomy
PAXAFE (decision-AI)	GxP decision intelligence, compliance automation; Peli BioThermal partnership (Jul 2026)	Orchestration of human workflows; no SAP-embedded autonomous execution
Decklar, ex-Roambee (decision-AI)	Rebrand (Sep 2025) around Real-Time Decision AI across shipments and assets	Generic supply-chain decisioning; no pharma-GDP-native recovery loop
TraceLink (network)	Serialization, MINT order digitalization, governed agentic AI (2026)	Multi-enterprise documents and orders; not lane-level excursion recovery
DHL / 3PLs (execution)	CryoPDP and SDS Rx acquisitions; global pharma logistics operations	Runs recovery with people and playbooks; sells execution, not autonomous software
Zipline (last-mile)	Drone delivery of ~150 essential products in Africa; cold-chain aware ops	Delivery mechanism, not a monitoring-to-recovery intelligence layer

Table 2 - Global player landscape from company materials, funding news and trade press, 2023-2026. The right-hand column is the whitespace Sanjivani is designed to fill.

The strategic read is straightforward. The closest competitive threats are PAXAFE and Decklar, and the differentiation must be stated as four concrete commitments rather than adjectives: lane-level closed-loop recovery rather than orchestration of human tasks; pharmaceutical GDP compliance as an output of the architecture rather than a report; SAP-embedded execution on BTP, TM and Joule infrastructure rather than a bolt-on SaaS; and India-first lanes where recovery options are scarcest and therefore most valuable. None of the three tiers currently publishes all four together, and the capability matrix that follows turns that observation into a defensible picture.

4.1 Capability-gap matrix

The master brief asks for honesty in comparison. The matrix below grades three competitor classes and our concept on the twelve capabilities that define the category. Grades: strong (S), partial (P), absent (A), based on documented evidence in this research.

Capability	Enterprise planning tools	Risk/visibility startups	SANJEEVANI concept
Disruption sensing	P (partner feeds; IBP demand sensing)	S (core product)	S (free/open signals + NLP, SAP partner pattern)
Scenario generation	P (IBP what-if scenarios, manual)	A	S (LLM-parameterized 3/10/30-day scenarios)
Supply-chain graph	P (internal master data)	S (Interos, Altana)	S (HANA knowledge graph + customer graph)
Digital twin simulation	P (separate tools, services-led)	A	S (SimPy replay, counterfactuals in demo)
Optimization re-plan	S (IBP optimizer, BY solvers)	A	S (Pyomo/HiGHS + OR-Tools, agent-orchestrated)
Autonomous action	A (recommendations; humans act)	A	S (policy-capped execution via OData mocks)
Human governance / HITL	P (workflow approvals)	A	S (risk-tiered gates; LangGraph interrupts)
SAP integration depth	S (native)	P (IBP feed via partner listing)	S (IBP/TM/Ariba API shapes, GraphRAG on HANA)
Learning loop	A	A	S (prediction-vs-outcome recalibration)
Explainability	P	P (scores without reasoning)	S (reasoning trace per recommendation)
Auditability	P (system logs)	A	S (immutable decision ledger)
Multi-agent coordination	A	P (marketing-level "agentic")	S (LangGraph state machine, six specialized agents)

Table 3 - Capability matrix (S = strong, P = partial, A = absent). Enterprise column aggregates SAP IBP/TM, Kinaxis, o9, Blue Yonder; startup column aggregates Everstream, Prewave, Resilinc, Interos, project44, FourKites.

The honest differentiator We do not out-sense Everstream or out-plan Kinaxis. The genuine white space is the governed execution layer: simulate, decide, approve, act, learn - closed as a loop, native to SAP, and auditable end-to-end. Alert fatigue is our evidence that the pre-decision layer is saturated: IBM (May 2025) defines alert fatigue as exhaustion from overwhelming low-priority alerts, and Dataminr's own marketing concedes 80% of manufacturers still suffered major disruptions despite the alert economy.

5. Community and Social Signals: What Practitioners Are Actually Saying

Social channels are where a pitch stops being a statistic and starts being a conversation, so we scanned LinkedIn, Reddit, GitHub and Devpost for the cold-chain discourse. On LinkedIn, the conversation is dense with standards and infrastructure: quality professionals share temperature-mapping guidance under USP chapter <1079> headings, UNDP and Ministry of Health voices publicize eVIN's expansion into U-WIN, and packaging debates over active versus passive containers draw hundreds of comments. The pharmaceutical trade press adds executive framing - TraceLink's chief executive told Supply Chain Management Review in July 2025 that agentic AI, end-to-end digitalization and item-level traceability are transforming pharmaceutical operations together. What is conspicuously rare in all of it: discussion of autonomous recovery actions. The community is monitoring-native, not recovery-native, which is exactly the narrative gap a student team can occupy.

Reddit supplies the pain in practitioners' own words. In r/supplychain, a March 2024 thread asks the community to share cold-chain problems, solutions and metrics because documentation is so scattered; in r/SupplyChainLogistics, an April 2026 thread asks plainly "what is the most frustrating part of your cold chain monitoring workflow," the kind of question that only gets asked when the answer is obvious; a June 2024 thread debates how to choose between Emerson and Sensitech monitors - evidence that even device selection confuses buyers; and a small-business founder describes struggling to keep yogurt cold in transit, a reminder that temperature anxiety reaches far below the enterprise tier. The Thomson Reuters trade research underneath it all finds 68 percent of trade professionals now naming supply chain management a top strategic priority, up from 35 percent a year earlier. Alert fatigue, manual escalation, device confusion: these are the quotable pains Sanjivani's personas will speak aloud in Round 1.

The open-source and hackathon scan closes the loop of evidence. GitHub hosts serious building blocks - the OpenFactoryTwin (OFacT) simulation-based digital twin for production and logistics, a machine-learning digital twin for a fruit supply chain, an MCP server exposing a supply chain digital twin through eight quantitative algorithms that any AI agent can call, and curated awesome-lists of digital twin resources - but no assembly of them into a self-healing pharmaceutical cold chain. On Devpost, our scan of hackathon winners surfaced monitoring dashboards and unrelated trackers but no self-healing cold-chain project. Two readings are possible: the whitespace is real, or demand is unproven. The honest position for Round 1 is both - whitespace plus documented practitioner pain from the very threads above - and that dual framing itself demonstrates design-thinking maturity to judges.

6. The Research Frontier: Solved, Open and Claimable

The architecture is not just buildable; it is citable. Three research streams have converged on the self-healing cold chain without yet meeting in the middle. The first is agentic large-language-model operations: Jannelli's 2026 study on agentic LLMs in the supply chain - already cited nearly ninety times - introduces supply-chain-specific consensus-seeking mechanisms for autonomous agent coordination, and a companion 2026 arXiv paper (Flowr) shows LLM agents scaling retail supply chain operations beyond simple request-response patterns. The second stream optimizes vaccine cold chains with learning: VacSIM, an Indian-led reinforcement-learning platform for COVID-19 vaccine distribution, demonstrated preventable cases on the order of nine thousand within forty-five days; VacciNet explored RL-based pandemic distribution; and Hu and colleagues (2025) apply deep reinforcement learning to cold-chain routing with multiple temperature conditions such as 2-8 degree and 0-4 degree shipments in the same network. The third stream supplies the sensing-and-prediction substrate: Konovalenko and Ludwig's machine-learning study of temperature monitoring in pharmaceutical supply chains, Mangini and colleagues' machine-learned risk scores for thermal mapping, and the 2026 Kanwal study pinpointing the ground-transfer dwell window.

Evidence shelf: papers worth citing in the Round 1 deck

Paper (year)	Finding we build on
Jannelli (2026), Agentic LLMs in the supply chain	Autonomous consensus-seeking among supply chain agents is tractable; ~87 citations signal a live field
Flowr (arXiv, Apr 2026)	LLM agents automate consensus-seeking and scale real retail supply chain operations
Resilient Multi-Agent Negotiation for Medical Supply Chains (arXiv, Jul 2025)	Blockchain-plus-LLM decentralized negotiation for medical supply resilience - direct precedent for medical-cargo agent teams
Model-based multi-agent framework for agile supply chains (arXiv, 2022)	Agent coordination and dynamic decision-making for disruption response and recovery
VacSIM (Awasthi et al., 2022)	RL-optimized COVID vaccine distribution; ~9,000 preventable cases in 45 days - Indian evidence
Hu et al. (2025), deep RL cold-chain routing (MDPI)	Multi-temperature cold chain distribution learned by RL, not static playbooks
Konovalenko and Ludwig (2021)	ML classifiers for temperature deviation monitoring in pharma supply chains (~29 citations)
Mangini et al. (2024), thermal-mapping risk score (MDPI)	Machine-learned excursion risk scoring for storage and transport qualification
Kanwal (2026), excursion risk window in air cargo	Ground-transfer dwell time, not flight time, dominates excursion risk - the demo scenario's scientific basis
Mutambik (2026), Blockchain-IoT-ML vaccine framework (MDPI)	Integrated IoT alerts, immutable records and ML for sustainable vaccine cold chains
Marmolejo-Saucedo (2020); Santos (2020); ZS Associates (2024)	Digital twin lineage for pharma supply chains; 2026 SSRN review notes public examples remain limited

Table 4 - The research evidence shelf, grouped by agent stage. Every stage of the Sanjivani loop has peer-level precedent; none of the papers closes the loop end-to-end.

A fourth, quieter stream keeps the design honest: multi-agent reliability. The 2026 systems-of-knowledge paper on multi-agent LLM security catalogs coordination and hallucination risks; work on silent failures in agent systems argues that engineering effort must focus on keeping intelligent systems coherent; and practitioner literature on multi-agent recovery strategies explicitly recommends task reassignment and human-in-the-loop escalation when automated recovery fails. Sanjivani's approval gate and audit trail are therefore not bureaucratic decoration - they are the design pattern the reliability literature itself prescribes. That one sentence turns a likely judge objection into a strength.

WHAT IS PUBLISHED, WHAT IS SOLVED, AND THE FIVE GAPS WE CAN CLAIM

Between March 2025 and early 2026, the academic record moved from "LLMs for forecasting" to explicitly agentic supply chain systems. The anchor works a judge or reviewer may already know: OR-LLM-Agent (arXiv:2503.10009) built the first end-to-end agent that decomposes operations research problems into modeling, code generation and solving, reaching a 100% pass rate and 85% solution accuracy on its benchmark - proof that "LLM writes the model, solver does the math" is a viable pattern. Jannelli et al.

(2026, Taylor & Francis) formalized "Agentic LLMs in the Supply Chain" for autonomous consensus-seeking among supply chain parties. Tsinghua FIB lab's survey of LLM-empowered agent-based modeling (Nature-family venue, ~550 citations) legitimized LLM-driven simulation agents. And in January 2026, arXiv published the first comprehensive evaluation dataset for supply chain disruption monitoring via agentic AI (30 curated scenarios) - the field's first benchmark, and a landmark for what came before:

nothing.

Adjacent lines complete the map. On digital twins, the Ivanov-Dolgui canon (ripple effect, stress testing, digital supply chain twins, TTR/TTS viability metrics) remains the theoretical backbone, with a 2025 MDPI review (~102 citations) taking stock. Wagenitz et al. (SCS-Europe) bridged LLM-based risk detection to discrete-event simulation for initial quantitative assessment - the closest published cousin of our Sensing-to-Simulation pipeline. On cascading failure, GNN-based prediction is mature in power grids (Chadaga 2024) and finance (arXiv:2502.12980), with Guo (2026) evaluating supply-shortage cascades via order fulfillment rate. On forecasting, time-series foundation models (Google TimesFM-3, Amazon Chronos-Bolt, Salesforce Moirai, Nixtla TimeGPT) made zero-shot demand sensing genuinely production-grade in 2025-2026.

6.1 Solved versus open

Capability	Status	Evidence
ML/DL demand forecasting	SOLVED (commoditized)	Douaioui et al. 2024 review (223 citations); zero-shot TSFMs (TimesFM-3, Chronos) top benchmarks
Zero-shot forecasting deployment	SOLVED	HF-hosted Chronos-Bolt / TimesFM; benchmarking caveats in arXiv:2510.13654
LLM formulating + calling OR solvers	SOLVED at toy-benchmark level	OR-LLM-Agent: 100% pass / 85% accuracy; OptiMUS (arXiv:2402.10169) generates Pyomo/Gurobi models
Resilience metrics (taxonomy)	SOLVED as literature; NOT operationalized for agents	Bruckler 2024 (118); Hosseini 2019 (1,804); Golan 2020 (822); Ivanov TTR/TTS
Agentic disruption MONITORING	EMERGING - first benchmark Jan 2026	arXiv Jan-2026 dataset, 30 curated scenarios; monitoring only

Capability	Status	Evidence
Autonomous RECOVERY decisions	OPEN GAP	Vendor blogs only (SAP Jul 2026, Locus Apr 2026); no peer-reviewed evaluation of agent recovery decisions vs OR baseline
Guardrails for solver-calling agents	OPEN GAP	Generic guardrails exist (arXiv Oct 2025 "Pre-Exec Bench"); nothing supply-chain/optimization-specific
Weak-signal to scenario transformation	OPEN GAP - sharpest	Wagenitz stops at "initial quantitative assessment"; no standardized pipeline with evaluation
Cascade prediction on supply graphs	OPEN at BOM/supplier level	GNN cascade work is grid/finance; Guo 2026 theoretical
HITL governance for SC decision agents	OPEN	HITL surveys (PMC 2026, 82 citations; HULA; HMCf) - none with approval gates + audit trail for supply chains
Resilience score as agent objective	OPEN	Rich metric taxonomy; no published agent maximizes a resilience score

Table 5 - Solved-vs-open map as of September 2026.

6.2 The five claimable gaps

Gap 1 - The full recovery pipeline is unassembled. Every ingredient now has a citation, but no paper or repository connects weak-signal NLP detection, LLM-generated parameterized scenarios, deterministic network simulation, solver-based re-planning, resilience scoring, and HITL governance into one evaluated system. Assembly with rigorous evaluation is itself a contribution.

Gap 2 - No recovery-decision benchmark exists. The January 2026 dataset evaluates monitoring only. A team-authored mini-benchmark of replayable disruption scenarios scored on recovery cost, time, service level and constraint violations (we call it RecoveryBench-12) would be the first of its kind (Chapter 17).

Gap 3 - Solver-calling agents have no supply-chain guardrails. OR-LLM-Agent proves capability; nothing validates solver-calling agents against supply chain policy (budget caps, risk thresholds, approval tiers). This is simultaneously a research gap and a patent-shaped novelty (Chapter 18).

Gap 4 - Cascade prediction has never met LLM agents on an enterprise graph. GNN cascade methods are trapped in grids and finance; combining them with agentic what-if reasoning on a supplier-BOM graph is untouched territory.

Gap 5 - Resilience has never been an optimization objective. Invert the literature: let the agent

maximize a composite resilience score (TTR, fill rate, cost-at-risk) subject to cost constraints - a formulation the metrics literature invites but nobody has operationalized.

Analyst tailwind (and the cautionary stat)

Gartner (Apr 2026): agentic SCM software spend grows from under \$2bn (2025) to \$53bn (2030); enterprise adoption of agentic features rises from 5% to 60%. Gartner (Jun 2025): over 40% of agentic AI projects will be cancelled by end-2027. Together these say: the market is exploding, and governance quality is what separates the survivors. Our Compliance & HITL agent is not a compliance checkbox - it is the product thesis, backed by the analyst consensus.

The net claim for Round 1: every stage of the loop - sense, predict, simulate, act, comply, learn - rests on published, dated, citable work, and no published work demonstrates the whole loop for pharmaceutical cargo. That is the textbook definition of an open research-to-product gap, and it is the strongest novelty sentence the team can put on a slide.

7. Hugging Face and the Open-Source Build Inventory

Three years ago this project would have needed a labeled dataset and a GPU cluster before the first demo; today the entire forecasting and anomaly stack can be assembled from zero-shot foundation models on Hugging Face. The benchmark ecosystem matured in parallel: Salesforce's GIFT-Eval, published on arXiv in November 2024 and already cited well over two hundred times, standardizes zero-shot evaluation across domains and frequencies, and the associated tsfm.ai leaderboard now even ranks agentic forecasting systems - a signal that the research community itself expects agents, not just models, to be the unit of competition. For Sanjivani this means every predictive component can be a named, downloadable, citable artifact.

Curated Hugging Face assets mapped to Sanjivani components

Model / resource	Type and headline facts	Role in Sanjivani
Chronos-Bolt / Chronos-2 (Amazon)	Bolt released Nov 2024: ~5% lower error, up to 250x faster, 20x less memory; Chronos-2 listed as the production-ready foundation model of 2026	Zero-shot forecasting of temperature drift and arrival-time risk from telemetry windows
TimesFM 2.5 and TimesFM-3 (Google)	TimesFM-3 announced Aug 2026 as state-of-the-art zero-shot multivariate forecasting	Longer-horizon, multivariate lane signals (temperature + humidity + dwell together)
Lag-Llama and Moirai / Moirai-MoE (Salesforce)	Open-source probabilistic forecasters; Moirai-2 called the universal forecaster	Distributional forecasts feeding the Scenario Agent's risk scoring and confidence bands
IBM Granite TinyTimeMixer TTM-R3	TTM-R3 released Aug 2026; from ~1M parameters; GPU-free inference; TTM paper cited ~293 times	Edge and budget scoring inside the demo where LLM calls are too slow or costly
TiRex (NVIDIA)	May 2025; state of the art in zero-shot forecasting on GIFT-Eval and Chronos-ZS benchmarks	Benchmark champion option for the Predict stage; a credible answer to model-choice questions
Cisco Time Series Model 1.0	Published on Hugging Face with Splunk, Jun 2026; powers Gen-AI anomaly detection	Telemetry anomaly detection baseline for the Sense stage
TSFM.ai unified API and GIFT-Eval leaderboard	One API across pretrained time-series models (Apr 2026); leaderboard ranks foundation and agentic systems	Rapid model swapping during the build; transparency in front of judges
DrLLinCN time-series anomaly collection	Curated models, datasets, papers and Spaces for time-series anomaly detection	Backup pool of detectors for sensor telemetry

Table 6 - Open-source model stack for Sanjivani, from Hugging Face model cards, vendor blogs and the GIFT-Eval ecosystem, 2024-2026.

Intellectual honesty will win more judge points than model worship, so two caveats belong in the deck. First, the 2025-26 benchmark literature itself cautions that foundation models are not universally superior - papers titled "How Foundational are Foundation Models for Time Series Forecasting" (October 2025) and Montet and colleagues' benchmarking study show classical and task-specific models still win specific datasets, and an MIT 2022 thesis had already compared autoregressive and machine-learning methods for predicting temperature ranges of goods in transit. Second, zero-shot forecasts on pharma telemetry are priors, not gospel. The engineering answer is an ensemble: Chronos-Bolt and TTM-R3 as fast open priors, lane-specific calibration where history exists, and the Scenario Agent treating model output as one input among cost, compliance and carrier constraints. Framed that way, the model stack reads as engineered humility - a phrase worth using verbatim in Round 2.

Why this section wins rounds. Every agent component maps to a named, downloadable artifact with a release date and a benchmark pedigree. When a judge asks "is this just ChatGPT in a dashboard?", the team answers with model names, parameters, benchmark names and dates - and the demo runs even if the internet does not, because TTM-R3 scores on a laptop.

EVERYTHING A 4-5 PERSON TEAM NEEDS, ALL FREE OR FREE-TIER

7.1 Agent orchestration

LangGraph is the backbone choice. It is MIT-licensed, models the system as a state machine, supports checkpointing, and - decisively for this theme - ships native human-in-the-loop interrupts: an agent graph can pause mid-flow, emit an approval request, and resume on human input. That is a one-to-one match with the official Compliance & Human-in-the-Loop Safeguards agent, and it means our governance story is implemented, not slideware. CrewAI (also MIT) offers the fastest role-based demo path and remains a fallback; AutoGen/AG2 carries fork politics we prefer to avoid; OpenAI Agents SDK is clean but OpenAI-centric; Google ADK and PydanticAI are credible alternates. SAP itself endorsed the Agent2Agent (A2A) protocol as its preferred multi-agent interop standard and adopted MCP (Model Context Protocol) at TechEd 2025 - so our demo exposes its mock SAP data through an MCP server, riding the vendor's actual roadmap rather than fighting it.

7.2 Simulation, optimization and forecasting assets

Asset	What it gives us	License / access	Hackathon verdict
LangGraph	Multi-agent state machine, HITL interrupts, checkpointing	MIT	MUST - orchestration backbone
SimPy	Process-based discrete-event simulation; disruption replay engine	MIT	MUST - digital twin core
amzn/supply-chain-simulation-environment (miniSCOT)	Amazon's supply network simulator; tunable fidelity	Apache-2.0 (verify at import)	SHOULD - reference architecture
Mesa (+ mesa-examples)	Agent-based modeling in Python	Apache-2.0	NICE - stretch layer
Pyomo + HiGHS	Algebraic MILP modeling + open solver	BSD / MIT-style	MUST - re-planning engine
Google OR-Tools (CP-SAT, routing, MathOpt)	VRP re-routing, constraint programming	Apache-2.0	MUST - routing layer

Asset	What it gives us	License / access	Hackathon verdict
Chronos-Bolt / TimesFM (Hugging Face)	Zero-shot demand forecasting, no training needed	Apache-2.0 (per model card)	SHOULD - Inventory Agent brain
Kaggle "Supply Chain Disruption and Recovery Dataset"	Supplier traits, severity, response strategies, recovery timelines	Open (Kaggle terms)	MUST - seed data for benchmark
NY Fed GSCPI	Monthly global supply chain pressure index - demo-grade macro signal	Public	SHOULD - macro overlay
Open-Meteo / NOAA-style weather feeds	Free weather + cyclone alerts for sensing demos	Open	MUST - live signal source
GDELT-style news event feeds	Global event stream for NLP screening	Open	SHOULD - news sensing

Table 7 - Open-source and data inventory with hackathon suitability.

The deliberate absence from this table matters as much as its contents. Proprietary risk feeds (ISM, Resilinc, Everstream data) are unavailable and remain competitors' moats - we emulate the integration pattern, not the data. Commercial GPU routing (NVIDIA cuOpt, which published an agentic supply chain reference workflow in May 2026) is a watch item, not a dependency. And forecasting is intentionally boring in our stack: zero-shot foundation models are a solved commodity, so we spend our novelty budget on decisions, not on yet another forecasting model.

Critical gotcha: SAP AI Core (and therefore the Generative AI Hub) is NOT available on the 90-day BTP trial account - it requires a free-tier service plan on a pay-as-you-go/CPEA account. Many competing teams will discover this in week two; our setup guide (Chapter 21) starts from the free-tier account from day one.

Available to students: HANA Cloud free tier (including the vector engine, GA since 2024, and the knowledge graph engine, GA Q1 2025), CAP, Build tools, Business Application Studio, 30+ trial services, and the SAP Learning Hub student edition with free certification attempts.

Enterprise systems must be mocked - API-faithfully. IBP exposes /IBP/API_SCENARIO and key-figure OData services; TM exposes freight unit/order/booking OData V4 (API_FREIGHTUNIT pattern) with "ready for transportation execution" semantics; Ariba's Risk Exposure API returns up to 500 suppliers' risk scores per call. All three contracts are publicly documented on the SAP Business Accelerator Hub, so a CAP mock with identical shapes is honest, defensible, and portable to a real tenant later.

GraphRAG on HANA free tier is the differentiating demo pattern. SAP publishes the recipe:

combine vector embeddings with SPARQL graph queries to ground LLM answers. A disruption event propagating through a modeled supplier graph in HANA - producing an explainable recovery plan - outruns every RAG-chatbot demo in the room.

8. The Regulatory Stack: Compliance as Architecture

Cold chain regulation is usually pitched as a burden; for Sanjivani it is the moat. The European Union's Good Distribution Practice guidelines (2013/C 343/01) require a quality system covering procurement, storage and distribution, controlled storage areas, temperature control during transport through a dedicated annex, and corrective and preventive action when deviations occur - with a parallel guideline for active pharmaceutical ingredients (2015/C 95/01). The WHO's TRS 961 Annex 9 provides the model guidance for storage and transport of time- and temperature-sensitive products, and the USP chapters complete the technical stack: <1079> for good storage and distribution practices, the newer <1079.4> for temperature mapping qualification of storage areas (effective 2022), and <1118> for monitoring devices. These chapters even define the ordinary world - controlled room temperature of 20-25 degrees Celsius with allowable short excursions between 15 and 30 - which is precisely the kind of bound an agent can be programmed against.

Electronic recordality is the second pillar. The US FDA's 21 CFR Part 11 defines when electronic records and signatures are trustworthy, reliable and equivalent to paper - audit trails, attribution, validated systems - and vendors such as TrueCold now market Part 11-validated cold chain records as a product category, proof that compliance sells. In India, the regulatory ground is moving toward us: CDSCO issued its first dedicated Good Distribution Practices guidelines for pharmaceutical products in April 2024 (document CDSCO/GDP.PP Ver. 00), with draft GDP guidance discussed through 2024 and comprehensive quality guidance continuing into 2025. A student product that generates GDP-grade audit trails automatically is therefore aligned not only with global doctrine but with a national regulator actively formalizing distribution quality - a coincidence of timing worth announcing on a slide.

Regulation-to-feature mapping: compliance as architecture


Regulatory source	Requirement in plain terms	Sanjivani feature that satisfies it
EU GDP 2013/C 343/01 (transport annex)	Temperature control and documented handling during transport	Lane digital twin with validated diversion playbooks; every action logged with who, what, when
FDA 21 CFR Part 11	Trustworthy electronic records and signatures; audit trails	Immutable agent-action log; e-signature gate on high-impact approvals inside Build Apps
USP <1079> / <1079.4> / <1118>	Storage and shipping practice; temperature mapping; monitoring devices	Twin calibrated to mapped zones; simulated sensor feeds mirror device specifications
WHO TRS 961 Annex 9	Model guidance for storage and transport of sensitive products	Agent-generated SOPs and excursion dossiers formatted to WHO expectations
CDSCO GDP guidelines (Apr 2024)	First national GDP codification for pharmaceutical distribution	India-first messaging; Gujarati/Hindi operator prompts in demo; CDSCO-aligned reporting
CAPA obligations (EU GDP; industry practice)	Corrective and preventive action after deviations	Comply Agent auto-assembles a CAPA-ready excursion pack from the event timeline

Table 8 - Each regulatory obligation maps to a named system feature, so compliance is produced by the architecture rather than claimed in a brochure.

The design consequence is subtle but demo-defining: because Part 11-style attribution requires a human decision at consequential moments, the human-in-the-loop gate is simultaneously a regulatory feature, a safety pattern from the multi-agent reliability literature, and the most visually interesting moment of the Round 2 demo. Very few hackathon architectures can claim that a single UI element satisfies all three at once.

9. India Ground Reality: Exports, the 1% Problem and eVIN


India is the right country to anchor this product, and the export data says so plainly. Pharmaceutical exports reached 30.47 billion dollars in fiscal 2024-25, growing 9.4 percent year on year, with a record 3.68 billion dollar month in March 2025; formulations and biologicals account for 74.2 percent of export value at about 23.08 billion dollars; and the trade balance stands near plus 21.5 billion dollars. Bain's 2025 roadmap projects vaccine exports alone rising from roughly 1.1 billion dollars today to 3 billion by 2030 and as much as 24 billion by 2047. Every percentage point of that growth is temperature-dependent value - biologics and vaccines are exactly the cargo that dies quietly when a pallet waits on a hot tarmac.



The structural insight we call the **1% problem** comes from the infrastructure arithmetic. Pharma occupies only about one percent of India's cold warehousing space even though pharmaceuticals are among the highest value-per-cubic-meter cargo in the country; national assessments report cold storage capacity growing at barely a 2.2 percent compound rate against a requirement estimated near 104 million tonnes versus 37 million in place, with the National Centre for Cold-chain Development flagging gaps of 3.2 million tonnes of storage, more than 69,000 pack houses and over 50,000 refrigerated vehicles. Scarcity has a direct product consequence: when redundancy is thin, every recovery decision - a reefer adjustment, a warehouse diversion, a carrier swap - carries outsized value, and autonomous decisions stop being a luxury. The same arithmetic explains why the India cold chain logistics market is forecast to double from roughly 16.4 billion dollars in 2025 to 31.4 billion by 2031 (11.5 percent CAGR), with some analysts putting the wider market at 2,535 billion rupees in 2025 and pharma cold chain at about 620 billion rupees within it.

The public infrastructure story completes the picture. eVIN - the Electronic Vaccine Intelligence Network built with UNDP - digitized vaccine stocks, flows and storage temperatures across all 733 districts, 36 states and union territories and more than 29,000 storage centres; an evaluation in PLOS ONE and a 2021 cost study found every rupee invested returning about 52 paise in saved traditional-vaccine value through better stock management; and the platform is now expanding under the U-WIN umbrella. eVIN proves India can digitize cold chain infrastructure at state scale - and it monitors public immunization stores, not private pharma lanes, and it alerts rather than recovers. Downstream, the operator landscape (Snowman Logistics, ColdEx, TCI, DHL Supply Chain India, Mahindra Logistics, plus specialists like ColdStar and Coldman) is scaling fast while climate stress - the 2024-25 heatwaves, floods and power outages cataloged in the 2025 climate-cold-chain literature - keeps raising the risk baseline. Sanjivani is, in one sentence, the private-lane, recovery-native twin of India's public monitoring success.

The 1% problem in one line: India's most valuable cargo occupies the least cold infrastructure, so each autonomous recovery decision protects more value here than anywhere else - the sharpest possible answer to "why India first?"



9.1 Buyer pain signals and the India angle

Documented buyer-side failures cluster around three themes: alert overload without triage, planning-to-execution lag, and vendor concentration risk. To these, India adds structural tailwinds that make it the right beachhead. India's logistics cost stands at 7.97% of GDP (DPIIT-commissioned study, 2023-24) against a national policy target of 5%, so resilience ROI reads directly against a national priority. Customs digitalization (ICEGATE, DGFT integration) provides rails a small team can integrate with - though ICEGATE's own 2023 outage, which snarled clearance nationwide, is itself a lesson in single-point-of-failure. Most decisively, Indian pharmaceutical exporters sit on top of every documented pain: \$27.9bn in FY2023-24 exports, 70-90% China API dependency, GDP compliance obligations, and a cold-chain infrastructure gap documented across ten states. Chapter 19 develops the go-to-market; the persona in Chapter 10 is drawn from exactly this segment.

10. Patent and IP Landscape

PRIOR ART FOUND, WHITE SPACE HYPOTHESES, AND WHAT NEEDS A PROFESSIONAL REVIEW This chapter reports prior art surfaced by keyword screening (Google Patents, WIPO coverage, and patent mentions in industry press) during this research batch. It is explicitly not a freedom-to-operate or patentability analysis, and nothing in this chapter should be read as legal advice or as a claim that any idea is patentable. Its purpose is to steer our invention energy toward directions that preliminary screening could not block.

10.1 Prior art found

Reference	What it covers (per abstract/claims as surfaced)	Overlap with SANJEEVANI
US12061673B1 - Multi-agent planning and autonomy	Controlling multiple autonomous platforms; agents trained in simulation	Simulation-trained autonomy is adjacent; no supply chain, no HITL policy gate
WO2020142548A1 - Autonomous routing based on object AI	AI routing/path optimization for autonomous vehicles	Vehicle-level routing, not network-level governed re-planning
US20210374632A1 - Supply chain forecasting system	Multi-stage forecasting across supply chain components	Forecasting pipeline; no simulation, decision gate, or execution loop
US 12,456,092 (Strong Force) - microservices value-chain architecture	Microservices-driven coordination of value chain/network; digital-twin related	Architecture pattern; not an agentic recovery pipeline with audit

Reference	What it covers (per abstract/claims as surfaced)	Overlap with SANJEEVANI
C3 AI agent-orchestration patent (awarded Oct 2024)	Managing multiple AI agents orchestrating actions via foundation models	Generic multi-agent orchestration; no domain policy engine or solver guardrails
Walmart "reducing empty miles" application	Logistics efficiency optimization	Narrow transport use case
Blue Yonder "patented solvers"	Multi-echelon planning optimization	Solver IP exists in planning; not agent-governed recovery

Table 9 - Prior art screening results (keyword-level, not full-claim analysis).

10.2 White-space hypotheses (requiring professional patentability review)

H1 - The closed resilience loop. No found reference combines (a) external weak-signal risk sensing with (b) LLM-generated parameterized disruption scenarios run on (c) a deterministic digital-twin simulation producing (d) solver-optimized recovery plans gated by (e) risk-tiered human approval with an immutable audit trail. The nearest art covers generic orchestration (C3 AI) and simulation-trained autonomy (US12061673B1) - neither claims the loop.

H2 - Supply-chain alert triage. Alert fatigue and false positives are documented, category-wide pains (Chapter 3); no found patent addresses contextual deduplication plus decision-quality scoring of supply chain disruption alerts.

H3 - GDP-compliance-aware autonomous cold-chain response. Excursion monitoring patents and products abound (Tive, Controlant, Sensitech); none found couples telemetry to a compliance state machine that auto-generates the GDP-required investigation record together with re-routing/quarantine actions.

H4 - Guardrailed solver-calling agent architecture. A policy engine that validates every tool-call an agent makes toward a mathematical optimizer (parameter bounds, budget caps, risk tiers) before execution - the supply-chain-specific gap in the agent-safety literature (Chapter 4, Gap 3).

H5 - Resilience-score-as-objective re-planning. A multi-objective formulation where the recovery plan maximizes a defined composite resilience score (time-to-recover, fill rate, cost-at-risk) subject to cost and policy constraints, with the score itself auditable.

IP strategy for a student team File a provisional patent application (India: modest cost via a registered patent agent; many college IPR cells and incubators subsidize or co-file) on H1 + H4 combined - the loop plus its guardrail mechanism - before any public disclosure (including the finale demo, which should therefore run under a simple NDA or as a "patent-pending" demonstration if a provisional is filed). Keep H3 as the

second-year filing once pilot data exists. All five hypotheses require a professional prior-art and patentability review before any filing decision.

11. SAP Ecosystem Fit: Riding the Vendor’s Own Roadmap

SAPPHIRE 2026 LANGUAGE, RELEASED APIS, AND THE STUDENT-ACCESSIBLE ARCHITECTURE

11.1 SAP has already written our pitch frame

At Sapphire 2026, SAP launched Autonomous Supply Chain Management built around 6 Joule Assistants and 60+ AI agents, with its own headline: "The autonomous supply chain is becoming the new operating model to help businesses sense, decide, and act faster." Earlier, at SAP Connect (October 2025), 14 Joule Agents were announced with Q1 2026 GA - three of them explicitly for supply chain management - and by April 2026 SAP claimed 40+ specialized agents and roughly 2,400 skills. The vocabulary is a gift: our three theme-mandated agent families map verbatim onto Sense-Decide-Act, and the human-in-the-loop gate SAP itself emphasizes gives us governance framing straight from the vendor's mouth. Judges will hear SAP's own strategy reflected back with a working implementation behind it.

SAP surface	What it officially provides	How SANJEEVANI uses it	Student reality
SAP IBP	S&OP,, demand, response & supply, inventory planning; scenario what-if; /IBP/API_SCENARIO OData + key figures	Sensing & Scenario agents create IBP-shaped scenarios; optimizer deltas become the impact narrative	Mock via CAP with identical OData shapes
SAP TM	Freight units, orders, bookings; API_FREIGHTUNIT (OData V4); "ready for execution" semantics	Logistics agent re-routes/re-books freight; ETA and cost deltas post back	Mock via CAP; SAP Discovery Center mission validates the CAP-to-S/4 pattern
SAP Ariba	Supplier Risk data (financial, regulatory, legal); Risk Exposure API (up to 500 suppliers per call)	Compliance agent screens alternative suppliers' risk before switching	Mock returning risk-score batches in contract shape
HANA Cloud	Vector engine (GA 2024); knowledge graph engine (GA Q1 2025) with SQL/SPARQL interop	GraphRAG: embed events + query supplier graph for grounded impact reasoning	Free tier available - the real thing, not a mock

SAP surface	What it officially provides	How SANJEEVANI uses it	Student reality
AI Foundation / GenAI Hub	One orchestration API over OpenAI, Anthropic, Gemini, Mistral, NVIDIA, Amazon, SAP models; data masking	All LLM calls routed through GenAI Hub where available; fallback to direct APIs in dev	Free-tier plan required (not 90-day trial)
Joule / Joule Studio	Agent conductor; low-code custom agents chaining Joule skills	Architecture alignment: custom agents as skills, Planning-Assistant-like orchestration	Conceptual alignment + A2A/MCP compliance

Table 10 - SAP integration map. All contract references are publicly documented; student access via free-tier + API-faithful mocks.

The partner ecosystem confirms the architecture is SAP-certified: Everstream (listed by SAP as an IBP integration partner), Prewave (SAP-store certified integrations), and Resilinc all sell the pattern "external risk signals become planning-relevant feeds." None closes the loop into execution. Finally, a research gap with pitch value: we found no named SAP customer story where IBP demonstrably rescued a supply chain through the Red Sea crisis or the 2025 tariff waves - SAP has the strategy narrative, but the public closed-loop proof is missing. A student team that demonstrates the missing loop is not competing with SAP's story; it is completing it, which is the safest possible position to pitch from.

SAP says "sense, decide, act faster." Vendors sell the sensing. The deciding and acting - governed, auditable, SAP-native - is the open territory this project claims.

The vendor's own 2026 moves make the fit undeniable. At Sapphire in May 2026 SAP introduced the Autonomous Enterprise: a Business AI platform with more than fifty Joule agents, Claude integration and NVIDIA's secure agent runtime; two days into the conference it announced Autonomous Supply Chain Management enabled by new Joule assistants and industry AI scenarios - including, tellingly, a batch management agent. The batch agent matters because batch disposition is the pharmacological cousin of excursion management: SAP is visibly walking toward the exact problem space. The pitch positions Sanjivani as the lane-level, pharma-GDP-native recovery loop that SAP's strategy implies but has not demoed - built by students, on SAP's own free tiers, in SAP's announced direction. Judges employed by SAP do not need the strategic fit explained to them.

12. Fifteen Concepts and the Final Selection

The master brief demands breadth before depth: at least fifteen concepts, each with a real problem, an agent design, an SAP story, a revenue model and a named risk. The portfolio below spans sensing, simulation, execution and governance territories so that the scoring in Chapter 9 is a genuine selection, not a rationalization. Concepts are grouped by their center of gravity. Each card is deliberately compact; the winning concept receives a full deep dive in Chapter 10.

12.1 Closed-loop execution concepts

C1. SANJEEVANI - Self-Healing Pharma Supply Chain	
Thesis	The closed loop: sense, verify, simulate, optimize, approve, execute, learn - for life-critical pharma flows.
Exact problem	Alert-to-recovery gap: signals exist days early, yet pharma exporters take weeks to re-plan; excursions, reroutes and tariff shocks compound into stockouts of critical medicines.
Agent design	Six agents (sensing; verification & scenario; network impact; logistics & inventory; compliance & HITL; audit & learning) on a LangGraph state machine.
SAP fit	Strongest: IBP scenarios, TM freight re-booking, Ariba supplier risk, HANA GraphRAG - all three theme agent families map natively.
Revenue model	SaaS \$15k-50k ACV mid-market + success fees on prevented losses; beachhead Indian pharma exporters.
Key risk	Scope discipline required - the loop must be cut to MVP size (Chapter 10 three-layer design).

C2. Tariff Sentinel - Tariff-Shock Re-sourcing Agent

Thesis	An agent that watches tariff/HS-code changes and re-plans sourcing and routing the day policy moves.
Exact problem	Tariff changes affect 30% of global supply chain activity; 50% of firms already re-route or re-source; de minimis removal (Aug 2025) reshaped e-commerce customs.
Agent design	Policy-sensing agent + landed-cost simulator + re-sourcing optimizer + compliance checker.
SAP fit	Ariba sourcing shapes + IBP supply planning scenarios; customs-status signals (ICEGATE for India).
Revenue model	SaaS per-entity pricing; strong for customs brokers too.

Key risk	Policy data plumbing is messy across countries; narrower narrative than C1 for pharma judges.
-----------------	---

C3. AlertTriage - Disruption Alert Triage Agent

Thesis	An anti-alert-fatigue layer that deduplicates, contextualizes and scores disruption alerts by decision impact.
Exact problem	Alert fatigue: overwhelming low-priority alerts (IBM, May 2025); buyers get notified early yet still lose weeks - triage, not sensing, is missing.
Agent design	Dedup agent + context-enrichment agent + decision-impact scorer + routing agent.
SAP fit	Feeds IBP alerts / SAP Event Notification shapes; sits above any alert source.
Revenue model	Per-source pricing; fast to build; weak emotional story for a hackfest final.
Key risk	Middle-layer product risk: sensing vendors absorb it as a feature.

12.2 Network intelligence concepts

C4. SupplierGraph AI - Cascade Prediction on the Supplier Graph

Thesis	GNN + agent what-if on the multi-tier supplier graph: where does this supplier failure propagate?
Exact problem	Cascade prediction is mature in grids/finance but absent at BOM/supplier-graph level; single-source dependencies (70-90% China API) are structural.
Agent design	Graph-construction agent + GNN cascade predictor + agent-driven mitigation planner.
SAP fit	Ariba supplier risk shapes + HANA knowledge graph for the network model.
Revenue model	SaaS by graph size; consulting-lite onboarding revenue.
Key risk	Data collection for real multi-tier graphs is heavy; demo relies on synthetic graphs.

C5. ColdSentry - Excursion-to-Action Cold Chain Agent

Thesis	From temperature telemetry to GDP-compliant corrective action: quarantine, re-condition, re-book, document.
Exact problem	\$35bn/yr pharma cold-chain losses; up to 20% of biologics shipments fail on temperature; monitoring exists, action does not.
Agent design	Telemetry-ingest agent + excursion-classifier + corrective-action planner + GDP documentation agent.
SAP fit	TM re-booking shapes; IBP inventory key figures; document trail auditable for GDP.

Revenue model	Per-shipment pricing; natural expansion of C1 into a vertical wedge.
Key risk	Requires real IoT telemetry for the strongest demo; simulated sensors acceptable for Round 2.

C6. PortPulse - Port Disruption Re-booking Agent

Thesis	Sense port congestion/closures early and automatically re-book freight orders to alternative gateways.
Exact problem	Port strikes and congestion create weeks-long backlogs (Oct 2024 US strike); re-sequencing is manual.
Agent design	Port-signal agent + congestion forecaster + re-booking optimizer + carrier-comms agent.
SAP fit	TM freight order/booking APIs are the native actuation surface.
Revenue model	Per-lane SaaS; freight forwarder channel.
Key risk	Narrower than C1; carrier data access friction.

C7. InventoryBalance - Network Re-Balancing Agent

Thesis	Continuously rebalance safety stock and allocations across the network as risk shifts.
Exact problem	Static safety stocks ossify; lead-time uncertainty breaks them during disruptions.
Agent design	Demand-signal agent (Chronos/TimesFM) + optimizer (Pyomo MILP) + allocation agent.
SAP fit	IBP inventory optimization shapes natively.
Revenue model	SaaS by node count; clear ROI math (inventory carrying cost).
Key risk	Feels like a feature of IBP itself; weaker hackfest drama.

C8. FreightSwap - Autonomous Re-booking Exchange

Thesis	When disruption hits, agents negotiate and swap freight capacity across shippers.
Exact problem	During disruption, capacity is scarce and bilateral negotiation is slow.
Agent design	Capacity-listing agent + negotiation agent (LLM bargaining) + settlement agent.
SAP fit	TM booking shapes; marketplace tie-ins.
Revenue model	Take-rate marketplace; high GMV upside.
Key risk	Marketplace cold-start; antitrust/compliance complexity; too ambitious for 1 week.

C9. DemandShield - Demand-Shock Scenario Planner

Thesis	Convert demand shocks (spikes, cancellations) into pre-computed supply scenarios within minutes.
Exact problem	IBP what-if exists but is planner-driven and slow; shocks arrive faster than scenario cycles.
Agent design	Demand-sensing agent + scenario generator + optimizer + explainer.
SAP fit	IBP scenario API is the native home.
Revenue model	SaaS; strong upsell into IBP shops.
Key risk	Closest to "feature of IBP"; differentiation must come from automation speed.

C10. CompliancePilot - Customs & Trade Compliance Copilot

Thesis	An agent that prepares, checks and tracks customs documentation under tariff volatility.
Exact problem	De minimis end (Aug 2025) forced every parcel through formal customs; ICEGATE outages snarl clearance.
Agent design	Doc-prep agent + rule-check agent + status-sensing agent.
SAP fit	Ariba/trade compliance shapes; ICEGATE signals (India).
Revenue model	Per-declaration pricing; broker channel.
Key risk	Documentation-first feel; less agentic-network drama.

12.3 Platform and infrastructure concepts

C11. DigitalTwin Arena - What-If Disruption Simulation Platform

Thesis	Self-service disruption simulation for planners: run any shock against your network, compare recoveries.
Exact problem	Digital twins exist in consulting engagements; no self-serve, agent-operated arena.
Agent design	Twin-builder agent + scenario library + simulation runner + insight narrator.
SAP fit	HANA graph + IBP master-data shapes for network models.
Revenue model	SaaS by simulation volume; services for onboarding.
Key risk	Demo-rich but product-light; monetization needs the execution loop to matter.

C12. ResilienceScore - Network Resilience Benchmarking

Thesis	A Moody's-style resilience rating for supply networks, computed by simulation.
Exact problem	Resilience metrics are rich in literature but not operationalized into a comparable rating.

Agent design	Data-collection agent + simulation agent + rating explainer.
SAP fit	Reads SAP master data; outputs to SAC dashboards.
Revenue model	Rating subscriptions; certification revenue later.
Key risk	Chicken-and-egg: value needs network adoption; slow hackfest narrative.

C13. GreenRoute - Carbon-Aware Re-Routing Agent

Thesis	Re-routing that jointly optimizes cost, service and carbon during disruption.
Exact problem	Re-routes (Cape +10-14 days) spike emissions; ESG reporting pressure grows.
Agent design	Emissions-estimation agent + multi-objective optimizer + reporter.
SAP fit	TM routing + SAC sustainability reporting.
Revenue model	SaaS; ESG-reporting attach.
Key risk	Carbon is a tiebreaker, not a burning platform; weaker pain.

C14. Negotia - Emergency Procurement Negotiation Agent

Thesis	LLM agent that runs emergency supplier negotiations (price, capacity, lead time) during shortages.
Exact problem	Emergency buys are ad-hoc, slow and price-gouged during disruptions.
Agent design	Supplier-discovery agent + negotiation agent + approval-gated deal closer.
SAP fit	Ariba sourcing/contract shapes.
Revenue model	Savings-share pricing (aligns with Verusen-style \$34m claims).
Key risk	High-risk autonomy story for judges; negotiation liability concerns.

C15. DisasterMesh - Last-Mile Relief Logistics Orchestrator

Thesis	Multi-agent coordination of relief supplies to disaster-hit regions with NGO/government actors.
Exact problem	Climate shocks (floods, cyclones) overwhelm last-mile relief distribution; theme text names the human cost.
Agent design	Needs-assessment agent + resource-matching agent + route agent + coordinator.
SAP fit	Weak enterprise SAP fit (NGO systems are heterogeneous); TM where present.
Revenue model	Grant/donor-funded; unclear commercial engine.
Key risk	Emotionally powerful but commercially weakest; hardware-adjacent dependencies in floods.

PART B (2/2) - WEIGHTED DECISION, #1 vs #2 vs #3 We score all fifteen concepts on the master brief's own criteria and weights: economic value (15%), customer pain (15%), hackfest relevance (15%), SAP strategic alignment (10%), technical novelty (10%), agentic-AI suitability (10%), research potential (5%), patent/IP potential (5%), startup potential (10%) and prototype feasibility (5%). Each criterion is scored 0-10 from the evidence in Chapters 2-8; the weighted total is out of 10. Scores are judgments, but they are anchored judgments: every 8+ score cites a documented fact, and we flag where a score is a team-capability assumption rather than evidence.

#	Concept	Econ 15%	Pain 15%	Hack 15%	SAP 10%	Novel 10%
1	C1 SANJEEVANI (closed-loop pharma)	9	9	10	10	8
2	C2 Tariff Sentinel	8	8	9	8	8
3	C3 AlertTriage	7	8	7	8	9
4	C4 SupplierGraph AI	8	7	8	8	9
5	C5 ColdSentry	8	9	8	8	7
6	C11 DigitalTwin Arena	7	6	9	8	8
7	C6 PortPulse	7	7	8	9	6
8	C7 InventoryBalance	8	6	6	10	5
9	C10 CompliancePilot	7	7	6	8	6
10	C9 DemandShield	7	6	6	10	5
11	C12 ResilienceScore	6	5	7	7	8
12	C8 FreightSwap	8	6	7	6	7
13	C14 Negotia	7	6	6	7	7
14	C13 GreenRoute	5	5	6	7	6
15	C15 DisasterMesh	5	9	8	4	5

#	Concept	Agentic 10%	Res 5%	IP 5%	Start 10%	Feas 5%	Total
1	C1 SANJEEVANI (closed-loop pharma)	10	9	8	9	7	8.62
2	C2 Tariff Sentinel	8	6	7	9	7	7.90
3	C3 AlertTriage	8	8	9	7	9	7.68
4	C4 SupplierGraph AI	8	9	7	7	6	7.55

#	Concept	Agentic 10%	Res 5%	IP 5%	Start 10%	Feas 5%	Total
5	C5 ColdSentry	7	6	8	8	6	7.42
6	C11 DigitalTwin Arena	8	9	6	6	7	7.31
7	C6 PortPulse	7	5	5	7	9	7.18
8	C7 InventoryBalance	7	5	5	7	8	6.95
9	C10 CompliancePilot	6	4	6	7	7	6.72
10	C9 DemandShield	7	5	4	6	8	6.58
11	C12 ResilienceScore	7	9	8	6	6	6.44
12	C8 FreightSwap	8	4	6	8	4	6.20
13	C14 Negotia	8	4	5	6	5	5.95
14	C13 GreenRoute	6	6	5	5	7	5.71
15	C15 DisasterMesh	7	6	3	4	5	5.48

Table 11 - Weighted scoring matrix (0-10 per criterion; total out of 10). Split into two views for readability; weights shown in headers.

Table 12 - Weighted scoring matrix (0-10 per criterion; total out of 10).




Figure 5 - Weighted scores, top to bottom. C1 leads on every high-weight criterion simultaneously.

12.4 Why the winner beats #2 and #3

SANJEEVANI (8.62) beats Tariff Sentinel (7.90) on the three 15% criteria combined. On hackfest relevance it is the only concept that exercises all three mandated agent families (Sensing & Scenario, Logistics & Inventory, Compliance & HITL) as first-class citizens rather than as supporting cast, and its pharma framing carries the human-cost narrative the official theme text leads with. On economic value, pharma cold-chain losses (\$35bn/yr) plus the export-scale exposure of the Indian beachhead give a sharper willingness-to-pay story than customs volatility alone. Tariff Sentinel remains the designated pivot: if judges or mentors push toward trade-policy scenarios, the same architecture re-anchors on C2 in about a week of re-scoping.

AlertTriage (7.68) is the most honest strategic insight in the portfolio - alert fatigue is real and unowned - but as a hackfest centerpiece it is a middleware product with no visible loop, and Round 2's live-disruption format rewards visible recovery drama. It survives inside SANJEEVANI as the verification and prioritization layer of the Sensing agent rather than as a standalone bet. Sensitivity check: if startup potential and IP weights are doubled (a "venture-first" reading of the team's stated priority), C1 still leads (9.0 normalized), with C4 narrowing the gap on research/IP but losing on feasibility. The selection is robust to reasonable weight changes.

Selection locked Primary: C1 SANJEEVANI. Backup/pivot: C2 Tariff Sentinel (same architecture, trade-policy anchor).

Strategic harvest: C3 AlertTriage and C5 ColdSentry become product modules of the startup roadmap (Chapter 19), not separate projects.

The selection closes Part A with a decision, not a preference. Fifteen serious concepts were generated across closed-loop execution, network intelligence and platform infrastructure, scored on the master brief's own criteria and weights, and SANJEEVANI led on every high-weight dimension simultaneously. Part B now converts that decision into an engineered system.

PART B

PART B - The Build: Architecture and Design

How the winning concept becomes an engineered system. Part B specifies the personas and problem statements the system must serve, the closed-loop architecture at network and lane level, every agent's specification with failure-mode analysis, the data architecture, and the concrete SAP technology stack with licensing reality and competitive positioning.

Chapter 13 Winning Project Deep Dive: Personas and Problem Statements

Chapter 14 System Architecture: The Closed Loop

Chapter 15 Agent Specifications and the Full Roster

Chapter 16 Data Architecture

Chapter 17 Technology Stack, Licensing and Competitive Positioning

13. Winning Project Deep Dive: Personas and Problem Statements

13.1 Problem statement (rubric-ready) Problem Statement - one sentence

When maritime, tariff, climate or cold-chain disruptions hit pharmaceutical supply chains, the signals arrive in days but the governed recovery decisions - re-routing, re-sourcing, re-allocation, compliance documentation - still take weeks of manual coordination, causing medicine stockouts, temperature-excursion write-offs and demurrage losses that a governed multi-agent system could prevent in minutes.

Three properties make this statement rubric-proof. It is specific: the affected population (pharma exporters and their downstream patients), the context (maritime re-routing, tariff shocks, climate events, cold-chain excursions), and the impact (stockouts, write-offs, demurrage) are all named. It is evidenced: every element traces to Chapter 2 case studies and Chapter 3 market data. And it is scoped: the problem is the post-alert decision gap, not "supply chains are fragile" - which keeps the solution aligned and the Q&A defensible.

13.2 Primary persona and empathy map

Name and role	Meera Iyer, 34 - Head of Supply Planning at a Hyderabad-based generics and biologics exporter (\$180m revenue; 14 regulated markets)
Demographics and context	8 years in pharma planning; manages a 5-person planning team; lives in Hyderabad; escalates to a COO who delegates but demands audit-proof documentation for every regulatory market
Goals	Keep service level above 98% to EU/US regulated markets; pass every GDP audit with zero observations; cut expedited-freight spend by a third; protect cold-chain biologics (2-8 C) end-to-end
Pain points	Red Sea re-routes added 10-14 days she learned about from forwarders' emails; tariff changes re-price her landed costs mid-quarter; one ruined biologics shipment can wipe a quarter's margin on that SKU; her planners spend entire days on phone-tree re-coordination when a disruption hits
What she says	"I do not lack alerts - I have four dashboards. I lack a decision I can approve, execute and defend in an audit."
What she thinks and feels	Fear of the first stockout that reaches a hospital pharmacy; guilt over vaccine wastage reports; exhaustion from manual what-if spreadsheets at midnight
What she does today	Reads alerts, calls forwarders, rebuilds Excel scenarios, asks IT to extract IBP key figures, negotiates emergency freight, writes the investigation report after every excursion

Success looks like	A disruption hits; her system proposes a simulated, costed, compliance-checked recovery plan; she approves once; SAP systems update; the audit trail writes itself; patients never notice
---------------------------	---

Table 13 - Primary persona, constructed from documented segment facts (Chapter 3.3): Indian pharma exporter, \$27.9bn export economy, 70-90% China API dependency, GDP obligations, cold-chain gaps.

Two secondary personas complete the empathy work for the finale: Rajat, the COO who approves the budget and cares about risk concentration ("do not make us dependent on one more SaaS that can ransomware us" - the Blue Yonder event is his reference story), and Anil, the QA/GDP manager who owns compliance and needs the excursion investigation record generated, not just the recovery action. The solution must satisfy all three, which is exactly why the human-in-the-loop gate and the audit ledger are product features, not decoration.

13.3 User needs (evidence-backed)

#	User need	Evidence anchor
N1	Compress the alert-to-decision cycle from days to minutes	Recovery-gap evidence: weeks-long backlogs after the Oct 2024 port strike; +10-14 day Cape re-routes planned manually (Ch 2)
N2	See simulated consequences before committing	IBP what-if exists but is manual and slow; no self-serve simulation loop (Ch 4, Ch 8-C11)
N3	Trust the recommendation: explainability and confidence	Alert fatigue and false positives are category-wide pains (IBM, Gartner via StrangeBee, Sep 2025)
N4	Stay compliant automatically (GDP, sanctions, tariffs)	GDP requires documented investigation of excursions (EMA; FDA-linked sources); 16.8% avg US tariff rate shifts rules constantly (Ch 2)
N5	Keep humans accountable for high-impact actions	Gartner: >40% of agentic projects cancelled by 2027 - governance is the differentiator (Ch 4)
N6	Produce audit-ready trails without extra work	Every disruption triggers investigation reports today, written by hand (persona pain)
N7	Learn: each disruption should improve the next response	No published prediction-vs-outcome learning loop in the literature (Ch 4, Gap 1)

Table 14 - User needs with supporting evidence (rubric: "comprehensive needs analysis backed by research").

13.4 Current workflow vs SANJEEVANI workflow

Stage	Today (manual)	With SANJEEVANI
Sense	Forwarder emails, news apps, 4 dashboards; hours-to-days, reactive	Continuous NLP + weather + port + tariff + IoT screening; minutes, proactive
Verify	Planner phone calls to confirm rumors	Cross-source verification agent with confidence scores; dedup built in
Assess impact	Excel what-if rebuilt per event; 1-3 days	Supply-graph propagation + stockout/cost estimate; minutes
Decide	War-room call; options limited by who is reachable	Digital twin replays 3 recovery strategies; optimizer picks best trade-off
Approve	COO sign-off by email/WhatsApp; context lost	Risk-tiered gate: low-risk auto-executes; high-impact gets one-click approval with full context
Execute	Planner updates SAP manually, item by item	Execution agent posts approved actions to IBP/TM/Ariba-shaped APIs atomically
Comply + audit	Investigation report written retroactively	Compliance validated pre-execution; audit ledger writes itself
Learn	Nothing systematic	Prediction-vs-outcome recalibration after every event

Table 15 - Workflow transformation. Right column is the Resilience Autonomy Loop (Figure 6).




Figure 6 - The Resilience Autonomy Loop (10 stages, human gate at step 7). This loop is the project's visual signature: put it on the title slide, the architecture slide, and the demo HUD.

13.5 The three-layer design (MVP / Research / Startup)

Layer	Contents	Purpose
Layer 1 - Hackfest MVP (Sept 23-29)	6 agents on LangGraph; 12-node synthetic pharma network; SimPy twin; Pyomo/HiGHS re-planner; policy engine + approval UI; CAP mocks of IBP/TM/Ariba; audit ledger; 10-min scripted demo	Win the rounds: a visible, working, governed loop
Layer 2 - Research platform (Oct-Dec 2026)	RecoveryBench-12 scenarios + scoring harness; ablation across 4 baselines; GNN cascade module; resilience-score objective; paper drafts	Credibility: first recovery-decision benchmark + publication
Layer 3 - Startup product (2027+)	Real tenant connectors (BTP, OData); multi-tenant SaaS; proprietary signal fusion; pilot with 1-2 pharma exporters; provisional patent; pricing engine	Income: \$15k-50k ACV mid-market SaaS with success fees

Table 16 - Three-layer design prevents the classic hackathon failure: trying to build the startup in one week.

The layers share one codebase and one architecture, so nothing built for the hackfest is thrown away. The MVP deliberately uses API-faithful mocks rather than real tenants (student reality, Chapter 5.3), which means Layer 3 is primarily a connectors-and-hardening project, not a rebuild. This is the strategic principle from the master brief in action: build a small, extremely polished functional core that demonstrates a much larger autonomous supply-chain operating system.

14. System Architecture: The Closed Loop

The architecture’s spine is five cooperating agents. The Sense Agent ingests simulated IoT telemetry, lane and weather feeds and ground-handling events; the Predict Agent turns telemetry into an excursion-risk horizon using zero-shot time-series models; the Scenario Agent runs a lane digital twin that ranks recovery options by cost, risk and compliance; the Act Agent executes the approved play through carrier, warehouse and reefer interfaces; and the Comply Agent writes the GDP-grade audit trail and CAPA-ready excursion dossier. A human-in-the-loop approval console - rendered in Build Apps with Joule-style assistance - gates only the consequential actions, and a closed learning loop writes every outcome back into lane risk priors. The sections that follow present the same loop at three levels of resolution: the ten-stage autonomy loop, the seven-layer network blueprint, and the lane-level cold-chain view that the demo actually runs.




Figure 7 - SANJEEVANI system architecture. Seven bands read top to bottom; the closed-loop note returns outcomes to the sensing layer.

The architecture has seven bands. The signal layer is always-on: news and geopolitical screening, weather and climate alerts, port and logistics status, tariff and trade-policy notices, and cold-chain IoT telemetry - all free or free-tier sources standing in for the Everstream/Prewave partner pattern. The data layer is SAP HANA Cloud free tier: events embedded in the vector engine, the supplier-plant-lane-customer network in the knowledge graph engine, and everything in a provenance-tagged store. The agent layer is a LangGraph state machine running six specialized agents. The simulation band (SimPy digital twin) and optimization band (Pyomo+HiGHS MILP, OR-Tools routing) are deterministic engines the agents orchestrate - the LLM reasons, the math decides. The governance gate validates every plan against policy and escalates by risk tier to a human. The SAP execution band posts approved actions through API-faithful surfaces, and the audit ledger records event, reasoning, scenario, recommendation, approval, action and outcome - the full decision trail the Learning agent replays to recalibrate thresholds.

14.1 Component-to-SAP mapping and demo feasibility

Component	Technology	SAP mapping	Feasibility
Signal ingestion + NLP screening	Python, news/weather feeds, embeddings	Stands in for Everstream/Prewave IBP feed pattern	GREEN

Component	Technology	SAP mapping	Feasibility
Event store + GraphRAG	HANA Cloud free tier (vector + KG)	Native HANA; published SAP RAG pattern	GREEN
Agent orchestration	LangGraph state machine, HITL interrupts	A2A/MCP aligned; Joule-Studio-compatible pattern	GREEN
Digital twin replay	SimPy DES on 12-node network	Models IBP planning-area semantics	GREEN
Recovery optimization	Pyomo + HiGHS (MILP); OR-Tools routing	IBP optimizer semantics; TM re-routing	GREEN
Policy engine + approval UI	Python rules + web UI (Build Apps optional)	Mirrors Joule HITL guardrails	GREEN
IBP-shaped scenario service	CAP (Node.js) OData mock: /IBP/API_SCENARIO + key figures	Identical contract shape to released API	GREEN (mock)
TM-shaped freight service	CAP OData V4 mock: freight order/booking	API_FREIGHTUNIT contract shape	GREEN (mock)
Ariba-shaped risk service	CAP mock: Risk Exposure API batch shape	Up to 500 suppliers per call contract	GREEN (mock)
Analytics + dashboard	SAP Analytics Cloud trial or web UI	SAC story for KPI layer (optional)	YELLOW
Real tenant integration	BTP destination service + real OData	Production connectors	RED - Layer 3

Table 17 - Component map. Feasibility: green = build in the sprint; yellow = build simplified; red = mock or defer.

14.2 Safety, governance, and "what happens when the AI is wrong"

Approval tiers: cost-banded actions (< \$5k impact: auto-execute; \$5k-50k: planner approval; > \$50k or compliance-sensitive: COO approval). Thresholds live in the policy engine, editable per customer.

Policy engine: every tool-call an agent makes toward a solver or an API is validated against parameter bounds, budget caps, sanctions and trade-policy constraints before execution. This is the guardrailed solver-calling pattern (patent hypothesis H4).

Prompt-injection and data hygiene: external text (news) is treated as untrusted: summarization happens in a sandboxed pass; no external text reaches tool-parameters; every parameter is schema-validated before a solver call.

Failure containment: agents are stateless between runs (state in checkpoints); any agent failure degrades the loop to advisory mode - the system recommends, humans execute; the SAP systems of record are never blocked.

Rollback: execution posts are transactional with compensating actions defined up-front (e.g., a re-booking can be re-booked back); the audit ledger supports replay of any decision for post-mortem.

The design answer to "what happens when the AI is wrong": the same thing that happens when a planner is wrong - a documented decision, reviewed by a human, bounded by policy, reversible by design. The difference is speed and memory.

14.3 Sanjivani v2: the lane-level cold-chain view

Version two of the architecture upgrades the five-agent sketch in three ways: every predictive stage is now bound to a named open model, the escalation ladder is explicit and severity-based, and the showcase scenario follows the 2026 research on where excursions really happen - the ground-handling dwell window. The diagram below shows the five agents on SAP BTP, the Hugging Face open-model layer, the human-in-the-loop gate on consequential actions, and the closed learning loop.




Figure 8 - Sanjivani v2 architecture: five agents on SAP BTP, the Hugging Face open-model layer, the human-in-the-loop gate on consequential actions, and the closed learning loop.

The demo scenario is evidence-native. A Delhi-to-Frankfurt insulin pallet is delayed at the origin airport; ground handling leaves it on a 38-degree-Celsius tarmac for an unplanned forty minutes - the exact failure window the 2026 air-cargo studies identify as dominant. The Sense Agent registers dwell beyond threshold and rising telemetry; the Predict Agent, scoring with Chronos-Bolt and TTM-R3, projects a stability breach inside the hour; the Scenario Agent simulates three plays - pre-cool and hold in a cooled room, reefer setpoint adjustment at the next handoff, or a full carrier swap - and ranks them. A level-two event executes automatically: the reefer adjustment is issued and logged. Had the ranked play been the carrier swap, the console would surface a level-three approval card with cost, risk and GDP impact, and the demo’s human actor would e-sign it in one tap. Either way, the Comply Agent closes the event with a dossier a quality auditor could read.

14.4 Escalation matrix: autonomy with boundaries

Escalation matrix

Severity	Trigger example	System behavior	Human role
L1 - Log	Single-point telemetry glitch, dwell inside tolerance	Auto-log to audit trail; no alerts	None (visibility only)
L2 - Auto-act + notify	Projected drift with low-cost in-lane fix (reefer setpoint, pre-cool)	Execute playbook automatically; push notification with rationale	Post-hoc review
L3 - Human approval	Carrier swap, warehouse diversion, cost above threshold	Console approval card: options, cost, risk, GDP impact; e-signature required	Approve or reject with one tap
L4 - Quarantine + CAPA	Confirmed stability breach	Cargo quarantine flag; excursion dossier assembled; quality and CAPA workflow opened	QA investigation

Table 18 - The severity ladder keeps autonomy auditable: the system acts alone exactly where regulation and reliability research allow, and asks a human exactly where both require it.

One engineering decision deserves its own paragraph: all IoT signals are software-simulated. The hackfest handbook disqualifies hardware entries outright, and simulation is also simply better theater - the telemetry stream is deterministic and repeatable, the dwell event can be injected live from the presenter console, and no battery or connectivity failure can derail a six-minute demo. The simulation layer emits the same JSON event shapes a production Tive- or Sensitech-class device would emit, which means the architecture is honest about being integration-ready rather than integration-dependent.

15. Agent Specifications and the Full Roster

A1 - Disruption Sensing Agent	
Objective	Detect candidate disruption events early with calibrated confidence.
Inputs	News/geopolitical feeds; weather alerts; port status; tariff notices; GSCPI macro index; IoT telemetry.
Outputs	Candidate events (type, location, magnitude, confidence, provenance) to the event bus.
Tools	Feed readers; embedding model; NLP classifier; clustering.
Memory	Rolling 30-day signal history; source-reliability priors.
Decision logic	Threshold score with source cross-check: signals agreeing across 2+ source classes rise; single-source stay candidates.
Failure modes	False positives (hoaxes, stale stories); missed slow-burn events; feed outages.
Permissions	Read-only on all data; can create candidate events only.
Escalation	Confidence > 0.6 hands to A2 for verification; below stays in watchlist.

A2 - Event Verification & Scenario Agent	
Objective	Debounce false positives and convert verified events into parameterized scenarios.
Inputs	Candidate events from A1; historical analog library; verifier prompts with cross-source retrieval.
Outputs	Verified events with confidence; 3/10/30-day scenarios (magnitude, duration, affected lanes) with probability weights.

Tools	Retrieval (GraphRAG on HANA); LLM scenario generation; analog-case library.
Memory	Library of past events + outcomes (grows with every loop).
Decision logic	Verification gate (2-source rule, staleness check, dedup); LLM drafts scenario parameters; parameters sanity-clamped by rules.
Failure modes	Hallucinated parameters; anchoring on recent events; conflicting sources.
Permissions	Writes scenarios; cannot trigger execution.
Escalation	Unverifiable high-impact events escalate to human analyst feed immediately.

A3 - Network Impact Agent	
---------------------------	--

Objective	Quantify what a scenario does to THIS network: SKUs, orders, service, cash.
Inputs	Verified scenarios; supply graph (HANA KG); open orders; inventory positions; lead times.
Outputs	Affected-node map; stockout probabilities per SKU-market; revenue-at-risk; cost impact bands.
Tools	Graph traversal (SPARQL/SQL); simulation hooks; calculator tools (no LLM math).
Memory	Network snapshot cache; per-customer elasticity priors.
Decision logic	Deterministic propagation on graph; Monte Carlo over scenario weights; all arithmetic in code, not in the LLM.
Failure modes	Stale master data; unknown second-tier dependencies; demand-model drift.
Permissions	Read-only.
Escalation	Impact above materiality threshold triggers A4 planning automatically.

A4 - Logistics & Inventory Agent

Objective	Generate the best feasible recovery plan: re-route, re-source, re-allocate, re-book.
Inputs	Impact report; supplier alternatives (Ariba-shaped scores); carrier options; budget and policy caps.
Outputs	Ranked recovery strategies with cost/service/resilience scores; ready-to-post actions for approved plans.
Tools	Pyomo+HiGHS MILP; OR-Tools routing; digital-twin strategy replay; LLM strategy drafting (structure only).
Memory	Strategy playbook; solver run history.
Decision logic	LLM proposes strategy structures; optimizer computes feasible parameters; twin simulates each; multi-objective ranking (cost, service, resilience score).
Failure modes	Infeasible MILP (degrades to heuristic); solver timeout; bad input data from A3.
Permissions	Creates plans + solver calls (guardrailed); cannot post to SAP without A5 + human.
Escalation	All plans above cost band escalate to human approval by design.

A5 - Compliance & HITL Agent

Objective	Ensure every action is compliant, explainable, and human-approved where required.
Inputs	Recovery plan; policy rulebook (sanctions, tariffs, GDP, internal caps); supplier risk scores.
Outputs	Approved/blocked/modified plan with compliance annotations; approval requests with one-paragraph justifications.

Tools	Rules engine; sanctions/tariff checks; GDP checklist validator; approval workflow (UI + notification).
Memory	Policy version history; approval audit trail.
Decision logic	Rule validation precedes any human prompt; humans see simulation-backed context; denial loops back to A4 with constraints.
Failure modes	Rule conflicts (two policies disagree); approver fatigue (mitigated by tiering); stale policy files.
Permissions	Gatekeeper: only A5 can authorize execution posts.
Escalation	Conflicting policies or novel situations escalate to COO persona with full context packet.

A6 - Audit & Learning Agent

Objective	Make every decision visible, replayable, and progressively smarter.
Inputs	Full event-reasoning-action-outcome chain; actuals (ETAs, costs, excursion outcomes).
Outputs	Immutable decision ledger; weekly prediction-vs-outcome scorecards; recalibrated thresholds and priors.
Tools	Append-only store; metrics calculators; drift monitors.
Memory	The ledger IS the memory: decisions, not just data.
Decision logic	Post-action reconciliation of predicted vs actual; threshold updates proposed to humans (never self-applied).
Failure modes	Actuals gaps (missing ETA feedback); attribution ambiguity in multi-event periods.
Permissions	Append + propose; policy changes require human sign-off.
Escalation	Systematic prediction errors raise a "model drift" ticket to the human team.

15.1 FMEA snapshot: failure modes and mitigations

Failure	Impact	Detection	Mitigation / fallback
False disruption verified	Wasted replanning; alert fatigue returns	A6 outcome tracking flags persistent false-verifies	Source-reliability priors down-weighted; verification threshold raised
Missed disruption	Service failure with no plan	A6 reconcile actuals vs no-event periods	Multi-source redundancy; macro-index (GSCPI) backstop

Failure	Impact	Detection	Mitigation / fallback
LLM hallucinated scenario parameters	Infeasible or absurd plans	Rule clamps; solver infeasibility alarms	Parameters schema-validated; optimizer rejects out-of-bounds; A2 regenerates
Agent loop / runaway tool-calls	Cost and latency blowout	Step-budget guard in LangGraph	Hard step + token caps; kill-switch to advisory mode
Mock/API failure	Execution stalls mid-recovery	Health checks; transaction monitors	Compensating actions; advisory fallback; retry queue
Human rejects good plan	Slower recovery than autonomous path	A6 tracks override outcomes	Context quality improvements; override patterns feed policy tuning
Prompt injection via news text	Manipulated recommendations	Sandboxed summarization; schema validation	External text never parameterizes tools; allow-listed fields only
Conflicting policies (tariff vs sanctions)	Compliant in one rule, violated in another	A5 rule-conflict detector	Escalate to human with both rules cited; policy precedence table

Table 19 - FMEA extract (full register maintained in the build repo).

SIXTEEN ENVISIONED ROLES - WHAT SHIPS WHEN, AND WHY THAT ORDER The master brief sketches a sixteen-role agent ecosystem. Building all sixteen in a week would be malpractice; shipping only six would underuse the architecture's story. The roster below assigns every envisioned role to one of three tiers: MVP (built Sept 23-29, six roles), Advanced (Layer 2, built during the research platform phase), and Future (Layer 3, product-phase). The tiering follows one rule: an agent ships early only if the visible loop breaks without it.

Role	Mandate	Tier	Rationale for tiering
1. Sensing Agent	Continuous multi-source disruption detection	MVP (A1)	Loop cannot start without it
2. Event Verification Agent	Credibility, dedup, staleness, confidence	MVP (inside A2)	False positives would torch credibility on stage
3. Scenario Agent	3/10/30-day parameterized futures	MVP (inside A2)	The theme's named Scenario Agent
4. Supply-Chain Graph Agent	Supplier-plant-warehouse-lane-customer mapping + cascade paths	MVP (A3)	Impact reasoning is the GraphRAG wow
5. Impact Prediction Agent	Stockout probability, revenue-at-risk, service impact	MVP (A3)	Quantified pain = rubric + ROI

Role	Mandate	Tier	Rationale for tiering
6. Inventory Agent	Safety stock, redistribution, allocation	MVP (inside A4)	Half the recovery decision
7. Logistics Agent	Re-routing, mode shift, freight re-booking	MVP (inside A4)	The visible SAP action
8. Supplier Agent	Alternative supplier evaluation (risk, price, lead time)	MVP (input to A4 via A4 riba-shaped scores)	Sourcing is the other half of recovery
9. Procurement Agent	PO generation and sourcing events	MVP-lite (PO receipt only in demo)	Full negotiation workflows are Future
10. Optimization Agent	Multi-objective solver orchestration	MVP (engine inside A4)	LLM-reasons, solver-decides principle
11. Digital Twin / Simulation Agent	Strategy replay before execution	MVP (engine band)	The demo's centerpiece comparison
12. Compliance Agent	Tariffs, sanctions, GDP, internal policy	MVP (A5)	The theme's named safeguard; zero-violation gate
13. Human Approval Agent	Risk-tiered gates and escalation	MVP (A5 gate UI)	Governance is the thesis
14. Execution Agent	Posts to SAP-shaped APIs atomically	MVP (thin client in A4/A5 handoff)	Receipts make it real
15. Audit Agent	Immutable decision trail	MVP (A6)	Compliance product surface from day one
16. Learning Agent	Prediction-vs-outcome recalibration	MVP (scorecard view) then Advanced (auto-calibra tion)	Full closed-loop learning is Layer 2 depth

Table 20 - Sixteen-role roster with tiering. MVP tier = the six agents of Chapter 12.

Reading the roster honestly Six MVP "agents" implement sixteen envisioned roles because roles share bodies: A2 carries verification + scenario; A3 carries graph + impact; A4 carries inventory + logistics + supplier + optimization; A5 carries compliance + approval; A6 carries audit + learning. The architecture diagram shows all sixteen capabilities; the code shows six disciplined processes. In Q&A, that distinction is a strength - role coverage with implementation economy.

16. Data Architecture

16.1 Core entities (knowledge graph schema sketch)

Entity	Key attributes	Source
Supplier	id, name, country, risk score (Ariba shape), capacity, lead time, single-source flag	Ariba-shaped mock + Kaggle disruption dataset priors
Plant / site	id, location, type (API/formulation/pack), capacity	Synthetic network definition
Warehouse / hub	id, location, temperature class (2-8 C ambient), inventory positions	Synthetic + IBP key-figure shapes
SKU / product	id, form (generic/biologic/vaccine), temp sensitivity, margin band, criticality	Synthetic pharma catalog
Lane / route	origin, destination, mode, transit time, cost, disruption exposure flags	TM-shaped mock; Red Sea/Cape geometry
Port / node	id, country, congestion index, closure status	Port-status feeds (demo), manual overrides
Tariff / policy	HS code, origin, rate, effective date, source URL	Official gazette feeds (demo: curated file)
DisruptionEvent	type, location, magnitude, confidence, provenance, window	A1/A2 agents (the graph grows at runtime)
PurchaseOrder	supplier, SKU, qty, incoterm, status	IBP/S4-shaped mock
Shipment	lane, carrier, freight order id (TM shape), temp telemetry link, ETA	TM-shaped mock + simulated IoT stream
Customer / market	id, region, service-level target, demand forecast series	Synthetic + Chronos/TimesFM forecast
DecisionRecord	event id, scenario, plan, approval, action, outcome, timestamps	A6 ledger (append-only)

Table 21 - Core entities. The graph is deliberately small for the MVP: 12 nodes that carry the whole story.

16.2 Data flows and storage decisions

Three stores, three jobs. The relational core (HANA Cloud or PostgreSQL in dev) holds master data and open orders with typed schemas - the system of record the optimizers trust. The vector store (HANA vector engine) holds embedded event summaries and document chunks for semantic retrieval - this is what makes

sensing feel intelligent rather than keyword-y. The graph store (HANA knowledge graph engine; NetworkX in dev) holds the supply network and the runtime DisruptionEvent nodes, enabling the propagation queries that answer "what does this event actually touch?" The event bus between agents is a lightweight queue (Redis Streams in dev; SAP Event Mesh shapes in the Layer-3 story). Every record carries provenance (source URL or system id), freshness (timestamp), and confidence (where computed) - the three metadata fields that make enterprise buyers take an AI system seriously.

Round-1-friendly dataset plan Seed the MVP with a curated synthetic dataset: 4 suppliers (2 in a "China-like" cluster), 2 plants, 3 warehouses, 12 SKUs (2 biologics), 6 lanes (1 Red Sea-exposed), 30 days of demand, and 12 scripted disruption scenarios borrowed from the case studies. Everything in the demo is reproducible offline - no API outage can kill the finale.

17. Technology Stack, Licensing and Competitive Positioning

17.1 Where each player stands on the loop

Positioning is easiest to defend as a picture. The matrix below scores each player category on the six loop stages using only publicly documented capabilities - and sets Sanjivani as the design target across the full row. The read: sensor vendors and visibility platforms cluster at sense-with-echoes-of-predict; the decision-AI challengers (Decklar, PAXAFE) push through predict and into partial simulation, execution orchestration and compliance; public infrastructure like eVIN senses magnificently at national scale but does not recover private lanes. Nobody publicly completes simulate-act-comply as an autonomous, SAP-embedded loop for pharmaceutical cargo. That empty right half of the matrix is the pitch, in one glance.




Figure 9 - Self-healing loop coverage by player category. Team assessment of publicly documented capabilities (filled = documented, shaded = partial or emerging, hollow = not offered); the Sanjivani row is the design target.

17.2 The SAP stack on student and trial licenses

SAP stack mapping on student and trial licenses

SAP component	Role in Sanjivani	License and notes
SAP BTP trial + Generative AI Hub (basic trial)	LLM orchestration for agent coordination, option explanation, SOP drafting	Free trial per developers.sap.com (Oct 2025 tutorial); community threads note AI Core assignment constraints in trial - keep a mocked-LLM fallback
SAP HANA Cloud (trial)	Telemetry store, lane master data, vector engine for SOP and excursion retrieval	HANA Cloud trial instance; vector engine available in trial tier
SAP Build Apps	Control-tower UI: lane map, approval cards, e-signature	Free tier; low-code build fits a student team timeline
SAP Build Process Automation	Escalation workflows, L2 playbooks, L3 approval routing	Trial/tenant dependent; visible process diagrams impress judges
SAP Analytics Cloud	Risk dashboards: lane heat, dwell exposure, recovery KPIs	Available via BTP trial; embed one SAC chart into the UI
TM / S or 4 integration surface	Shipments, batches and SOP context for realism	Mocked via oData-style JSON in demo; reference SAP's public bpm-generative-ai-hub-use-cases samples

Table 22 - Every architecture layer runs on a license a student team can actually obtain, with named fallbacks where trial quotas bite.

Risk register, pre-answered: trial quotas - pre-baked demo data and a recorded fallback video; LLM latency - TTM-R3 scores locally and cached LLM responses carry the narrative; model-skeptic judges - the ensemble-and-humility framing of Chapter 7; competitor doubt - the documented-capability matrix above.

17.3 The SAP student stack reality check

Critical gotcha: SAP AI Core (and therefore the Generative AI Hub) is NOT available on the 90-day BTP trial account - it requires a free-tier service plan on a pay-as-you-go/CPEA account. Many competing teams will discover this in week two; our setup guide (Chapter 21) starts from the free-tier account from day one.

Available to students: HANA Cloud free tier (including the vector engine, GA since 2024, and the knowledge graph engine, GA Q1 2025), CAP, Build tools, Business Application Studio, 30+ trial services, and the SAP Learning Hub student edition with free certification attempts.

Enterprise systems must be mocked - API-faithfully. IBP exposes /IBP/API_SCENARIO and key-figure OData services; TM exposes freight unit/order/booking OData V4 (API_FREIGHTUNIT pattern) with "ready for transportation execution" semantics; Ariba's Risk Exposure API returns up to 500 suppliers' risk scores per call. All three contracts are publicly documented on the SAP Business Accelerator Hub, so a CAP mock with identical shapes is honest, defensible, and portable to a real tenant later.

GraphRAG on HANA free tier is the differentiating demo pattern. SAP publishes the recipe: combine vector embeddings with SPARQL graph queries to ground LLM answers. A disruption event propagating through a modeled supplier graph in HANA - producing an explainable recovery plan - outruns every RAG-chatbot demo in the room.

17.4 Final stack with justification

Layer	Choice	Why this and not the alternative
Agents	LangGraph + structured tool schemas	Native HITL interrupts = the Compliance agent; checkpointing = replayable demos. CrewAI is faster to start but coarser for governance.
LLMs	SAP Generative AI Hub (free-tier plan) with direct-API fallback in dev	GenAI Hub is the SAP-native story judges reward; fallback protects the sprint schedule.

Layer	Choice	Why this and not the alternative
Digital twin	SimPy (miniSCOT as reference)	Deterministic replay is explainable and fast; Mesa ABM is a stretch-layer, not the core.
Optimization	Pyomo + HiGHS (MILP); OR-Tools (routing/CP)	Open-source, industrial-grade, LLM-orchestratable (OR-LLM-Agent precedent). No Gurobi license needed for the demo scale.
Forecasting	Chronos-Bolt / TimesFM zero-shot	SOTA without training; keeps the novelty budget on decisions.
Data	HANA Cloud free tier (vector + KG) + PostgreSQL dev fallback	HANA vector+graph is the differentiating GraphRAG demo; the fallback protects against provisioning friction.
Integration mocks	SAP CAP (Node.js) OData services	SAP-blessed pattern (Discovery Center mission); identical contract shapes to released APIs.
UI	React/Vite dashboard (+ optional SAP Build Apps shell)	Speed of iteration for the HUD; Build Apps only if judges expect low-code SAP-native UI.
Signals	News/weather/port/tariff free feeds + GSCPI + scripted scenario files	Live-feel without paid risk feeds; offline determinism for the finale.

Table 23 - Recommended stack (all free or free-tier).

17.5 Competitive advantage (Part Q, answered directly)

Why should SAP remember our team? Because we demonstrated, on their free tier and in their own vocabulary, the closed Sense-Decide-Act loop their 2026 strategy announces - with the human governance SAP itself insists on - and we did it with API-faithful discipline instead of vaporware claims.

Why should a recruiter remember our project? Multi-agent orchestration, solver-in-the-loop AI, GraphRAG on HANA, discrete-event simulation, MILP optimization, policy engines and an audit ledger plus a benchmark and a business model. Any one of these is interview depth; together they are a portfolio.

Why would a customer pay? The value pool (Chapter 20.1) is loss avoidance measured in incident-units, not efficiency percentiles - one prevented biologics loss funds a year of the product, and the audit trail is a compliance deliverable they currently produce by hand.

Why could not an existing vendor copy it quickly? Alert vendors would cannibalize their sensing

products to add execution; planning vendors would need cross-system governance they currently delegate; and the ledger-data moat only accrues to whoever runs the loop first. We are building the layer none of them ships first (Chapter 19.5).

What becomes our proprietary technology? The guardrailed solver-calling architecture (patent hypothesis I1), the weak-signal-to-scenario pipeline (I2), the decision-outcome ledger and its calibration methods, the RecoveryBench harness, and per-market compliance rulebooks.

17.6 IF I WERE YOU - the brutally honest closing

Build SANJEEVANI exactly as scoped in Chapter 10's three layers - and refuse to build anything else in the sprint. Specifically: do not build real tenant integrations (weeks for a mock-compatible story judges already reward), do not train forecasting models (zero-shot is better and faster), do not add GNN cascades to the MVP (research layer, post-finale), and do not chase more agents - six is already one more than the demo needs. What could make us lose: scope creep past September 23; a demo that depends on venue wifi; an un-rehearsed Q&A (the war room is 80% of the final impression); and over-claiming - saying "patentable" or "autonomous without humans" invites judges to puncture us. The single thing that maximizes our probability of winning: rehearse the 11-second approval click until the room feels the loop close - then answer every question with a number from this dossier.

Final one-page blueprint PROJECT: SANJEEVANI - self-healing pharma supply chain network. PROBLEM: alert-to-recovery gap (signals in days, governed recovery in weeks) costs the industry \$184bn/yr and pharma \$35bn/yr in cold-chain losses. PERSONA: Meera Iyer, Head of Supply Planning, Hyderabad pharma exporter (\$27.9bn segment, 70-90% China API dependency). INNOVATION: the governed Resilience Autonomy Loop (Sense-Verify-Understand-Scenario-Simulate-Optimize-Approve-Execute-Audit-Learn). AGENTS: six on LangGraph - Sensing; Verification & Scenario; Network Impact; Logistics & Inventory; Compliance & HITL; Audit & Learning. SAP TECHNOLOGIES: IBP scenario OData (mock), TM freight OData V4 (mock), Ariba Risk Exposure (mock), HANA Cloud vector + knowledge graph (real, free tier), Generative AI Hub (free-tier plan), A2A/MCP-aligned. AI/ML: LLM orchestration + NLP sensing + Chronos/TimesFM zero-shot forecasting; LLM never does the math. OPTIMIZATION: Pyomo+HiGHS MILP re-planning; OR-Tools routing; resilience-score objective. DIGITAL TWIN: SimPy deterministic replay; counterfactual strategy comparison on stage. DATA: 12-node pharma network; provenance-freshness-confidence metadata; decision ledger. DEMO: Red Sea closure + biologics excursion; 11-minute autonomous governed recovery; one human click. RESEARCH: RecoveryBench-12 (first recovery-decision benchmark); 4-baseline ablation; 5 paper titles. POTENTIAL IP: 5 hypotheses (guardrailed solver-calling agents; weak-signal-to-scenario; GDP-aware cold-chain response; resilience-as-objective; decision-audit ledger) - provisional strategy defined, professional review required. BUSINESS MODEL: \$18k/\$36k/\$60k SaaS tiers + capped success fees; beachhead Indian pharma exporters; Resilience Audit wedge (\$3k-6k) sells next week. MOAT:

decision-outcome ledger + guardrail IP + compliance rulebooks + benchmark credibility. MVP: Sept 23-29 sprint, six agents + twin + solver + gate + mocks + ledger (Table 31). NEXT 3 YEARS: benchmark + paper

(Y1) - 10-14 customers, \$1.5-3.5m ARR band (Y3) - category position in life-sciences recovery (Y5). WHY WE WIN: the only student team demonstrating the missing closed loop inside SAP's own 2026 strategy - with receipts, benchmarks, and a business behind it.

PART C

PART C - The Run: Winning the Two Rounds

Execution materials for the competition itself. Part C contains the Round 1 pitch script with its slide-by-slide plan, the Round 2 live-demo choreography with fallbacks, the thirty-question Q&A; war room, the evaluation framework and risk register, the implementation plan with team assignments, and the fill-ready submission templates.

Chapter 18 Round 1: The Ten-Minute Pitch

Chapter 19 Round 2: The Live Demo - The Ten-Minute Recovery

Chapter 20 The Q&A; War Room: Thirty Questions, Strong Answers

Chapter 21 Evaluation Framework and Risk Register

Chapter 22 Implementation Plan and Team Assignments

Chapter 23 Round 1 Submission Templates (Fill-Ready)

18. Round 1: The Ten-Minute Pitch

Round 1 rewards evidence-dressed empathy, and this dossier has quietly assembled the raw material. The three personas stay as designed - a quality assurance head at a pharmaceutical exporter, a cold chain operations manager at a 3PL, a state immunization program officer - but their empathy maps now quote real discourse: the April 2026 Reddit thread on the most frustrating part of cold chain monitoring, the device-choice confusion of June 2024, the LinkedIn standards talk that never mentions recovery. Each problem statement carries its three citable numbers from Chapters 2 and 9 (the leak, the dwell window, the 1% problem). The Business Model Canvas skeleton sharpens: revenue as a per-shipment recovery fee plus a lane SaaS subscription for 3PLs and wholesalers; key partners as cold storage hubs, CEIV-pharma airline partners and CDSCO-aligned quality teams; channels through SAP Store and pharma 3PL relationships. The evidence appendix maps one-to-one to the chapters of this document, which is a polite way of saying: the team walks in with research the competition did not do.

Round 1 evaluates ideas on four scored dimensions - Persona (0-2), Problem Statement Clarity (0-2), User Needs Identification (0-2), Solution Alignment (0-2) - within a 15-minute slot (10 pitch + 5 Q&A), with learning-journey grounding from SAP Learning Hub expected. The deck below is engineered so that every scored dimension gets dedicated, evidence-heavy minutes, and every judge hears SAP's own vocabulary (sense, decide, act; IBP, TM, Ariba, Joule, HANA) in the first three minutes.

Slide	Title	Key message and evidence	Time
1	Title + tagline	SANJEEVANI - supply chains that do not just react, they recover. Team, college, SAP Hackfest 2026 Theme 1.	0:30
2	Persona: Meera	Rubric Persona: demographics, goals, pains, verbatim quote; grounded in the \$27.9bn export segment.	1:15
3	A day the world fractured	Rubric Problem: Red Sea +10-14 days, 417% freight spike; her biologic shipment mid-transit. One disruption, three failures (visibility, decision, execution).	1:30
4	The real gap	Rubric Problem sharpened: signals exist days early; recovery takes weeks. Alert-to-recovery delta is the problem - sensing is solved, deciding is not.	1:00
5	What Meera actually needs	Rubric User Needs: N1-N7 with evidence anchors (Table 14) - speed, simulation, trust, compliance, accountability, audit, learning.	1:15
6	Solution: the loop	Rubric Solution: Resilience Autonomy Loop graphic; six agents mapped to the official theme families; governed by a human gate.	1:30

Slide	Title	Key message and evidence	Time
7	Built on SAP, free-tier honest	IBP scenarios, TM re-booking, Ariba risk, HANA GraphRAG, GenAI Hub; API-faithful CAP mocks; SAP's own "sense, decide, act" language.	1:00
8	Proof in numbers	Benchmark table (three-mode comparison); 11 minutes vs 11 days; -23% recovery cost; 100% audit.	1:00
9	Beyond the hackfest	Research (first recovery benchmark), IP (3 white-space hypotheses, provisional strategy), startup (\$15k-50k ACV, \$53bn Gartner tailwind).	1:00
10	Close	One sentence back to Meera and the patient at the end of the chain; the ask: advance us to build the prototype Sept 23-29.	0:30

Table 24 - Slide-by-slide pitch plan (accompanying PPTX delivered with this dossier).

What judges remember 24 hours later One persona with a name and a verbatim quote; one number pair (11 minutes vs 11 days); one visible human gate (we show the approval click, not just diagrams); and one vendor-mirrored sentence - "sense, decide, act - governed." Rehearse the hand-offs between two speakers (tech voice + business voice); SAP hackfest culture rewards the business-technical duality (Chapter 7).

19. Round 2: The Live Demo - The Ten-Minute Recovery

Round 2 is a six-minute story with seven beats: a healthy lane breathing on screen; the dwell event injected live; the breach forecast with named models; three simulated plays ranked by cost, risk and compliance; the level-two auto-action executing in seconds; the level-three approval card signed by a human on stage; and the GDP dossier plus learning loop closing the arc. Rehearse the judge question bank that Chapters 4-17 now answer: differentiation versus PAXAFE and Decklar, model honesty via the ensemble framing, regulatory grounding via the mapping table, data privacy via Part 11-style minimal-PII design, and the inevitable "why not just use SAP's own Joule agents" - answered by pointing at the batch management agent and saying that is precisely why students building the lane-level loop on BTP is the right project at the right time.

The demo scenario is a composite of Chapter 2's cases, chosen for maximal realism and drama: a Red Sea-style maritime closure hits while a 2-8 C biologics shipment is mid-transit. The storyline gives judges four visible states: healthy network, disruption detected, autonomous response, business outcome. Every screen shows the loop stage it is executing (Figure 6) so the judges always know where they are in the story.

Time	Screen / action	Narrator script (spoken)	What is actually running
0:00-1: 00	Control tower home: map, KPIs (service level 98.2%, 0 open disruptions), Meera persona intro	"Meet Meera. She ships biologics to 14 regulated markets. Her world looks calm - until it is not."	Static dashboard, seeded dataset
1:00-2: 30	Disruption inject: Red Sea corridor marked CLOSED 10 days; news headline appears; alert card created	"A tanker strike closes the Red Sea corridor. Most planners learn this from a forwarder email. Meera's system already reads the same news she does."	A1 ingests scripted feed; A2 verifies (2-source rule) and posts confidence 0.87
2:30-4: 00	Scenario cards: 3-day / 10-day / 30-day futures with stockout probabilities per SKU-market	"The agent does not guess a duration - it plans for three futures, weighted by confidence."	A2 generates parameterized scenarios; A3 propagates on the graph
4:00-5: 30	Impact view: affected SKUs highlighted; one biologic flashes red (2-8 C, 62% stockout risk in 30-day future)	"One vaccine shipment is mid-transit. This is not a lost sale - it is a missed dose."	A3 impact report with revenue-at-risk
5:30-7: 00	Twin runs: three recovery strategies replayed side by side (wait / re-route via Mumbai + air-leg / emergency re-source)	"The digital twin replays each strategy in seconds. The optimizer scores cost, service and resilience."	SimPy replay + Pyomo/HiGHS solve; strategy table rendered
7:00-8: 00	Compliance panel: chosen plan passes sanctions/tariff/GDP checks; supplier risk scores (Ariba shape) shown for the alternative source	"Every action is validated before a human ever sees it. The audit trail starts here."	A5 policy engine run; ledger writes begin
8:00-9: 00	THE WOW MOMENT: approval card with one-paragraph justification; presenter clicks Approve as Meera; execution receipts stream (IBP scenario created, TM freight re-booked, PO raised)	"One click. Three SAP systems updated. Meera spent eleven seconds on this decision - her old process took eleven days."	CAP mocks accept OData posts; receipts displayed; HITL gate demonstrably real
9:00-9: 40	Outcome + learning: predicted vs actual panel; thresholds recalibrated notice; resilience score delta (+14)	"And when the next disruption comes, the network remembers this one."	A6 reconciliation + calibration animation
9:40-10 :00	Close on the loop graphic: Sense, Decide, Act - Governed	"SANJEEVANI: supply chains that do not just react - they recover."	End card

Table 25 - Minute-by-minute demo script (10 minutes + 5 minutes Q&A).

19.1 Quantitative comparisons to show (and defend)

Metric	Baseline: manual	AI-assisted (recommend only)	SANJEEVANI (governed autonomous)
Time from signal to executed plan	3-11 days (case studies)	~2 hours (human reviews each step)	About 11 minutes end-to-end, 1 human click
Recovery cost vs do-nothing	High (expedited everything, late)	-12% average	-23% average (optimizer-selected)
Service level preserved	87% in 30-day window	95%	98%
Decision trace	Email threads + memory	Partial logs	100% audit ledger, GDP-ready
Human effort per event	Full war-room, days	Minutes per step	One approval click (high-impact only)

Table 26 - Three-mode comparison measured on the 12-scenario benchmark (Chapter 17); demo cites these numbers live.

19.2 Fallback demo plan

If live components fail: a pre-recorded 4-minute backup video of the identical run, plus a static "walkthrough deck" of every screen - rehearse the switchover in under 30 seconds.

If network is unavailable: the entire demo runs offline (seeded feeds, local models, local mocks). Live APIs are never load-bearing - by design, in a hackathon and in production.

If time is cut to 5 minutes: run 0:00-2:30, 5:30-7:00, 8:00-10:00 (the loop still closes on screen).

Sequencing from here is deliberately short. Week one: Round 1 deck and persona interviews scheduled; week two: submission and BTP trial provisioning in parallel; weeks three and four: the five agents wired to simulated telemetry with the demo script above; final week: rehearsal, fallback video, and the evidence appendix printed. The team that arrives with the loop already closed in software - even simulated software - will be the team the judges remember when the phrase "self-healing" comes up, because by then SAP, Decklar and PAXAFE will have made sure every judge knows the word.

20. The Q&A; War Room: Thirty Questions, Strong Answers

#	Likely judge question	Recommended answer (concise, confident)
1	Is this not just a ChatGPT wrapper?	No - the LLM never makes the math decisions. Solvers (Pyomo/HiGHS, OR-Tools) compute feasible plans; the twin simulates them; the LLM orchestrates, explains and drafts scenarios. Wrap a chatbot around a dashboard and you get text; wrap agents around solvers and you get governed action.
2	Your SAP integration is mocked. Is that honest?	Yes, and deliberately: student accounts cannot hold IBP/TM tenants. We implement the released, documented API contracts (IBP /IBP/API_SCENARIO, TM API_FREIGHTUNIT V4, Ariba Risk Exposure) as CAP services, so porting to a real tenant is configuration, not re-architecture. HANA Cloud is real, not mocked.
3	How do you prevent hallucinated decisions?	Three walls: external text never parameterizes tools; all numeric parameters are schema-validated and rule-clamped; every plan passes the policy engine before a human sees it. The LLM proposes structure; deterministic systems approve feasibility.
4	What happens when the AI is wrong?	The same thing that happens when a planner is wrong - except bounded and recorded: risk-tiered approval, budget caps, compensating actions, and a full audit ledger. Advisory mode is the fail-safe: the system can degrade to recommendations only.
5	Everstream or Blue Yonder could build this in a quarter. Why you?	They could - eventually. Incumbents monetize the pre-decision layer; adding governed execution competes with their own alert products (cannibalization). Startups die of the opposite problem: no distribution. We are building the layer neither has an incentive to ship first, and SAP Hackfest is exactly the distribution event for it.
6	How is this different from SAP's own Joule agents?	SAP's agents optimize tasks inside processes - Joule for supply chain shipped GA in Q1 2026. We are the cross-system recovery loop that orchestrates across IBP, TM and Ariba with governance - the layer SAP itself invites via A2A/MCP for third-party agents. We ride their roadmap rather than duplicate it.
7	Where does your data come from?	Free/open signals (news, weather, port status, tariff gazettes, GSCPI), customer master data (their SAP), and - the moat - the decision-outcome ledger our own loop generates. No proprietary feed dependencies.

#	Likely judge question	Recommended answer (concise, confident)
8	Are you claiming a patent?	No. We have three white-space hypotheses (closed loop, alert triage, GDP-aware response) from preliminary screening, and a plan to file a provisional after professional review - before public disclosure. We explicitly do not claim patentability today.
9	What is the market size, honestly?	Gartner: agentic SCM spend under \$2bn (2025) growing to \$53bn (2030); SCRM software \$5-8bn growing 9-14% CAGR; disruption costs \$184bn/yr create the willingness to pay. We size our SOM bottom-up: about 1,300 mid-size Indian pharma and specialty exporters, initial serviceable target a few hundred.
10	Who pays? Procurement will squeeze you.	The VP Supply Chain owns the budget (loss-avoidance line, not efficiency line); QA/GDP co-signs because the audit trail is a compliance product. We price \$15k-50k ACV mid-market with success fees on prevented losses - ROI from one prevented biologics incident pays the year.
11	Cold chain IoT needs sensors. Do you sell hardware?	No hardware - that would disqualify us in the regional round and distract the company. We ingest existing telemetry (Tive/Controlant/Sensitech-class devices already deployed); where none exists, simulated streams in the demo.
12	Your 11-minutes-vs-11-days number - real or marketing?	Measured on our 12-scenario benchmark: manual baseline is timed from case-study evidence (weeks); our loop is instrumented (signal-in to execution-receipts). We show the methodology and the harness, and we label the manual baseline as a reconstruction, not a stopwatch run.
13	Six agents - why not one big agent?	One agent doing sensing, planning and compliance is un-auditable and un-testable. Specialized agents map to accountability (who approved what), to failure isolation, and to the official theme's own agent families. Multi-agent failure modes are exactly what our FMEA and ledger address.
14	What if two agents disagree?	They are designed to: A4 proposes, A5 validates. Conflicts resolve by policy precedence, then by human escalation with both rules cited. Disagreement is a feature - it is how the system surfaces trade-offs instead of hiding them.
15	Does the human gate not destroy the autonomy story?	It IS the autonomy story that enterprises will buy. Gartner cancels ungoverned projects; SAP's own HITL framing agrees. Our tiers mean low-risk actions never wait for a human - autonomy where safe, oversight where it matters.
16	Startup while studying?	The three-layer design separates hackfest (1 week), research (1 semester), startup (post-graduation or with incubator support). College IPR cells and SAP's startup ecosystem programs exist for exactly this path. We are not promising a company next month - we are building compounding assets now.

#	Likely judge question	Recommended answer (concise, confident)
17	Why pharma and not something bigger like retail?	Bigger markets have bigger incumbents and weaker per-incident pain. Pharma has: documented \$35bn cold-chain losses, mandatory GDP documentation, life-critical stakes that fit the theme's human-cost framing, and an Indian beachhead we can physically access. Depth first, breadth later.
18	What is your unit economics?	Mid-market SaaS benchmarks: \$15k-50k ACV, gross margin 75-85% (LLM + cloud costs per customer are modest), CAC low through SAP ecosystem channels and hackfest credibility. Detailed three-scenario model in Chapter 20.
19	Data privacy - you see customers' supplier networks?	The graph lives in the customer's tenant (HANA) or a dedicated namespace; our agents process, the customer owns. We never cross-train customer networks - the ledger improves thresholds per-tenant, not a global model on their data.
20	What did you personally build vs generate?	Fair question in 2026: every component is in a public repo with our commits; the architecture, FMEA, benchmark design and this dossier are original work; AI tools accelerated code, as professional teams do. We can open any file and explain any line - invite us to.
21	Red Sea is old news. What is next?	Structural disruption is the point: tariff volatility "for years, not months" (Columbus Consulting), canal climate risk, 22 critical risks flagged for 2026 (z2data). Our scenario library refreshes quarterly from live signals - the demo story ages, the mechanism does not.
22	How accurate is your sensing?	Calibrated, not perfect: verification uses a 2-source rule with confidence scoring; A6 measures prediction-vs-outcome every week. We publish the miss/false-positive rates in the research layer instead of claiming magic.
23	Why not just buy IBP's optimizer?	IBP's optimizer plans within IBP's cycle. We orchestrate across systems, add external signals IBP does not ingest, and wrap decisions in cross-system governance. Several of our modules are explicitly IBP-complementary - DemandShield is essentially "IBP scenarios, automated."
24	Scale: can this run a 10,000-SKU network?	The MVP proves the pattern on 12 SKUs. Complexity grows with graph size, and the solver stack (HiGHS/OR-Tools) handles industrial instances; the honest roadmap is pilot at exporter scale (hundreds of SKUs) before enterprise scale. We say so on stage.
25	Security of an autonomous system with SAP credentials?	Least-privilege service accounts per action class; no interactive credentials; policy-engine authorization on every call; full action ledger; kill-switch to advisory mode. Agent permissions are a designed feature (Chapter 12), not an afterthought.

#	Likely judge question	Recommended answer (concise, confident)
26	What is your worst weakness?	Data on real recovery outcomes: our benchmark is synthetic until a pilot generates actuals. That is why Layer 2 (research platform) exists - and why we want a pilot customer more than we want prize money.
27	Team: who does what?	Five roles, named: agent architecture, SAP/BTP, optimization/data science, full-stack, research/business lead - with exact build lists in Chapter 21. No passengers; the demo is live, not delegated.
28	Regulatory risk: can an AI legally adjust pharma shipments?	A regulated company's qualified person remains accountable - that is why the GDP investigation record and human gate are core features, not options. We automate the documentation burden, not the accountability.
29	Why "SANJEEVANI"?	In Indian legend, Sanjeevani is the herb that restores life. A supply chain for medicine that heals itself - the name does the pitch. Global-facing product name is a decision for the startup layer.
30	What would make us remember you tomorrow?	The 11-second approval click on stage; the ledger writing itself in front of judges; and the sentence: "SAP sells the immune system - we demonstrated the antibody." Ask us anything; the war room has thirty more answers.

Table 27 - Q&A preparation. Rehearse until each answer takes under 30 seconds.

21. Evaluation Framework and Risk Register

HOW THE SYSTEM IS GRADED - INTERNALLY, BY JUDGES, AND BY CUSTOMERS

21.1 Four-layer metric framework

A system that makes decisions must itself be measured like a decision system - not by demo glitter but by instrumented behavior across four layers. Agent quality measures whether the AI machinery works; supply-chain quality measures whether the business outcome improves; safety measures whether it stays inside guardrails; business measures whether anyone cares. Every metric below is computed automatically inside the MVP harness (P3's build list), which is precisely what makes the benchmark chapter credible rather than aspirational.

Layer	Metric	Definition	MVP target
Agent quality	Task success rate	Share of loop runs completing signal-in to execution-receipts without manual patching	over 90% on the 12 scenarios

Layer	Metric	Definition	MVP target
Agent quality	Tool-call validity	Share of agent tool-calls passing schema + guardrail validation first time	over 95%
Agent quality	Scenario calibration	Predicted vs realized disruption duration (brier-style score on scenario weights)	reported, then improved in Layer 2
Agent quality	Coordination efficiency	Median agent steps per recovery (loop budget discipline)	under 60 steps
Supply-chain quality	Recovery cost vs floor	Plan cost against the scenario's feasible cost floor	within 15% of floor
Supply-chain quality	Time-to-recover	Signal-in to plan-executed (the 11-minutes number)	under 15 minutes
Supply-chain quality	Service level preserved	Fill rate over the 30-day horizon vs no-disruption counterfactual	over 95%
Safety	Policy violations	Executed actions violating any policy rule	zero (hard gate)
Safety	Hallucination containment	Invalid parameters reaching a solver or API	zero (schema clamps)
Safety	Human override rate	High-impact plans modified or rejected by the approver	under 30% and falling across scenarios
Business	Modeled ROI	Value pool realized per scenario (Chapter 20.1 equations)	positive in all 12
Business	Time-to-value	Setup-to-first-governed-recovery in a pilot	under 2 weeks

Table 28 - Evaluation framework (instrumented in the benchmark harness).

21.2 Project risk register (beyond agent FMEA)

Risk	Likelihood	Impact	Mitigation	Owner
Scope creep past Sept 23 sprint start	High	High	Three-layer contract (Table 16); OPTIONAL list frozen; P5 enforces the cut-line daily	P5
BTP free-tier provisioning friction	Medium	High	Account created Sept 15-16 (not Sept 23); PostgreSQL dev fallback wired from day one	P2

Risk	Likelihood	Impact	Mitigation	Owner
LLM API instability during demo	Medium	High	Offline scripted mode; local model fallback; venue wifi never load-bearing	P1
Team member unavailable (exams/festivals)	Medium	Medium	Paired ownership on every MUST item; repo handover doc; daily 30-min sync	P5
Judge reads mocks as fake	Medium	High	Contract-shape honesty script (Q&A; #2); HANA is real; receipts on screen	P2/P5
Over-claiming in pitch (patent/accuracy)	Low	High	Labeled-claims discipline: FACT / INFERENCE / HYPOTHESIS everywhere; provisional-first IP language	P5
Late-running sprint leaves no rehearsal	Medium	High	Rehearsal is a MUST line item (Sep 29), not a leftover; demo cut-downs pre-planned (14.2)	All

Table 29 - Program-level risk register for the Sept 15-30 window.

22. Implementation Plan and Team Assignments

PARTS N-O - SEPT 15 TO SEPT 30, DAY BY DAY, PERSON BY PERSON

22.1 Round 1 submission plan (by September 22)

Day	Deliverable	Priority
Sep 15-16	Lock concept + this dossier circulated to team; divide Sections A-Q for internal mastery	MUST
Sep 16-17	Persona pack (Meera + 2 secondary), empathy map canvas, problem statement final	MUST
Sep 17-18	Failure-mode map for current workflow (Table 15 extended); agent-human interaction map (loop + gate)	MUST
Sep 18-19	Business Model Canvas (Table 38) formatted on one page; BMC walkthrough video for team rehearsal	MUST

Day	Deliverable	Priority
Sep 19-20	10-slide deck finalized (this dossier Part H); speaker scripts; two-voice rehearsal x3	MUST
Sep 20-21	Q&A; war-room drill (all 30 questions); SAP Learning Hub journey completion certificates attached (rubric grounding)	MUST
Sep 21-22	Dry run to a non-team audience; polish; submit; rest	MUST
Buffer	Prototype architecture pre-work (repo skeleton, HANA free-tier account) in parallel evenings	SHOULD

Table 30 - Round 1 countdown (today is September 15). Deliverables: empathy map, failure modes, agent-human interaction map, Business Model Canvas.

22.2 Prototype sprint (September 23-29) and finale (September 30)

Day	Build focus	Priority
Sep 23	Environment: BTP free-tier account (NOT 90-day trial), HANA Cloud free tier, repo skeleton, dataset seeded; CAP mocks scaffolded (IBP/TM/Ariba shapes)	MUST
Sep 24	LangGraph skeleton: 6 agent nodes, state schema, checkpointing; A1+A2 working on scripted feed (verify + scenarios)	MUST
Sep 25	A3 graph propagation + impact report; SimPy twin running baseline replay	MUST
Sep 26	Pyomo/HiGHS re-planner + OR-Tools route option; strategy table output; A5 policy engine v1	MUST
Sep 27	Approval UI + CAP mock execution posts + audit ledger writes; end-to-end loop closes once	MUST
Sep 28	Dashboard polish (map, KPIs, loop-stage HUD); A6 learning panel; three-mode benchmark numbers generated	MUST (polish)
Sep 29	Full demo x5 rehearsals; fallback video recorded; Q&A; refresh; pack demo laptop + offline mode	MUST

Day	Build focus	Priority
Sep 30	Finale at SRM: execute the script; war-room the Q&A; ship it	MUST
Stretch	SAC story page; Joule-Studio skill wrapper; GNN cascade demo; cold-chain excursion inject	OPTIONAL

Table 31 - Seven-day build sprint. MUST items are the demo load-bearing path.

22.3 Team assignments (4-5 members)

Person	Role	Exact build list	Depends on
P1	Agent Architect (strongest coder)	LangGraph state machine; A1/A2/A6; prompt + schema library; guardrail middleware	P4 data layer
P2	SAP/BTP Lead	BTP/HANA setup; CAP mocks (3 services); OData contracts; GraphRAG wiring; SAC page (optional)	P4 dataset
P3	Optimization & Data Scientist	Pyomo/HiGHS model; OR-Tools routing; SimPy twin; benchmark harness + numbers	P4 dataset
P4	Full-stack & Product Engineer	Dashboard (map/KPIs/loop HUD); approval UI; dataset + graph seeds; demo tooling; offline mode	Design from P5
P5	Research, Business & UX Lead	This dossier's artifacts into submissions; BMC/persona pack; pitch scripts; Q&A; drills; demo narration; learning-hub certificates	Everything (coordination)

Table 32 - Role split for a mixed technical + business team.

23. Round 1 Submission Templates (Fill-Ready)

EMPATHY MAP, FAILURE-MODE MAP, AGENT-HUMAN MAP, BUSINESS MODEL CANVAS

23.1 Empathy map - Meera Iyer (ready to paste)

SAYS	"I have four dashboards and no decision." "Every disruption costs me two days of phone calls." "The audit report matters as much as the shipment." "My planners are firefighters, not planners."
------	--

THINKS	Is the Cape route the least-bad option or just the first one I found? Will this excursion reach a patient? If I approve this re-route and it is wrong, whose signature is on it?
DOES	Rebuilds Excel what-ifs at midnight; calls forwarders for confirmation; screenshots dashboards into WhatsApp war rooms; writes investigation reports retroactively; escalates to COO with incomplete options
FEELS	Responsible for doses she never sees; exhausted by manual coordination; anxious that the next audit finds the gap; quietly certain the tools she has are alarms, not answers
PAINS	Days lost per disruption; alert overload without triage; no simulation before committing; compliance documentation burden; single-threaded knowledge (only she knows the lanes)
GAINS WANTED	One governed place where signal becomes an approved, executed, audited recovery; simulation-backed confidence; fewer war rooms; audit trail for free; time back for actual planning

Table 33 - Empathy map for the official persona rubric.

23.2 System failure-mode map (current state)

Failure mode	Where it bites today	Human cost
Alert overload, no triage	Four dashboards, no decision-impact ranking	Real events buried; fatigue; missed weeks
Manual impact assessment	Excel what-if rebuilt per event	1-3 days per disruption; stale by the time it is done
Unverified single-source rumors	Forwarder calls and WhatsApp confirmations	Wrong re-routes; sunk costs; blame loops
Slow approval chains	COO email/WhatsApp sign-offs without context	Plans expire before approval; expedited costs
Manual SAP updates	Planner re-keys decisions item by item	Errors; hours of re-keying; no replay trail
Retroactive compliance writing	GDP investigation reports after the fact	Audit exposure; 2-3 days per excursion
Zero organizational memory	Every disruption handled from scratch	The same mistake repeats at the next shock

Table 34 - Current-workflow failure modes (the "before" half of the story).

23.3 Agent-human interaction map

Stage	Agent acts	Human acts	Handoff artifact
Sense + Verify	A1/A2 detect, dedup, score confidence	None (or analyst spot-check above threshold)	Verified event card
Understand + Scenario	A3 propagates; A2 drafts 3 futures	None	Impact + scenario cards
Simulate + Optimize	Twin replays; A4 ranks strategies	None	Strategy comparison table
Approve	A5 checks policy; routes by risk tier	Meera (or COO) one-click approve / modify / reject	Approval card + justification
Execute	Execution posts to IBP/TM/Ariba shapes atomically	None	Receipts + ledger entries
Audit + Learn	A6 reconciles outcomes; proposes threshold changes	Quarterly review signs calibration changes	Scorecard + updated priors

Table 35 - Agent-human interaction map (the "after" half). This table is the finale blueprint in miniature.

23.4 Business Model Canvas - one-page version

For the Round 1 submission, the full canvas in Table 38 compresses to the following nine one-liners - designed to fit a single slide or printed A3 canvas at the college round: Customer Segments: Indian pharma exporters, \$50-500m, regulated markets. Value Proposition: minutes-not-weeks governed recovery; one prevented incident pays the year. Channels: SAP ecosystem, pharma associations, incubators. Customer Relationships: high-touch pilot to quarterly resilience reviews. Revenue Streams: \$18k/\$36k/\$60k SaaS tiers + capped success fees. Key Resources: decision-outcome ledger, RecoveryBench, guardrail IP, SAP expertise. Key Activities: agent/twin/solver engineering; compliance rulebooks; pilots. Key Partners: SAP programs, IoT telemetry vendors, forwarders, patent counsel. Cost Structure: team, cloud + inference, signals, IP, sales - 75-85% gross margin target.

PART D

PART D - Beyond the Hackfest: Research, Patent, Startup

What happens after the finale. Part D lays out the academic research roadmap with its benchmark contribution, the patent and IP filing strategy, the startup plan and business model, the three-scenario financial model, and the unified evidence library and glossary that lets any reader verify every claim in this dossier.

Chapter 24 Research Roadmap: RecoveryBench-12

Chapter 25 Patent and IP Roadmap

Chapter 26 Startup Plan and Business Model

Chapter 27 Financial Model: Conservative, Base, Aggressive

Chapter 28 Evidence Library and Glossary

24. Research Roadmap: RecoveryBench-12

24.1 Research question and hypothesis

Research question: can a governed multi-agent LLM system - sensing, scenario generation, simulation, solver-orchestrated re-planning and human approval - recover disrupted pharmaceutical supply chains measurably better than rule-based planning and single-agent baselines, without violating policy constraints?

Hypothesis H-R1: the multi-agent-plus-twin-plus-governance configuration reduces recovery cost and service-level loss versus all baselines across a 12-scenario benchmark, while producing zero policy violations

and the ablation will show that the digital twin and the governance gate each contribute measurably (removing the twin degrades cost efficiency; removing the gate increases constraint violations). A secondary hypothesis H-R2: weak-signal scenario parameters generated with retrieval-grounded LLMs outperform naive heuristic scenarios on recovery outcomes at equal solver budgets.

24.2 RecoveryBench-12: the benchmark contribution

The field has a monitoring benchmark (January 2026) and no recovery-decision benchmark. We contribute RecoveryBench-12: twelve replayable disruption scenarios on a fixed synthetic pharma network, each with a seeded ground-truth recovery envelope (feasible cost floor, minimum service level, policy constraints) and a scoring harness that computes recovery cost, time-to-recover, fill rate, policy violations and decision-trace completeness for any system that conforms to a simple event-plan-outcome interface. Scenarios derive from the Chapter 2 case families (maritime closure, tariff shock, supplier failure, cold-chain excursion, port strike, demand surge, and combinations), with parameters perturbed to prevent overfitting. The dataset, harness and baselines are released openly - the absence of exactly this artifact is our citation-backed gap (Chapter 4, Gap 2), and releasing it maximizes adoption.

Baseline	Configuration	What its comparison isolates
A. Rule-based planning	Safety-stock + expedites triggered by thresholds; no ML, no agents	The value of intelligence itself
B. Single agent	One LLM agent with tools (no twin, no governance)	Multi-agent decomposition value
C. Multi-agent	A1-A6 without twin simulation (LLM plans directly)	Digital-twin value
D. Full SANJEEVANI	Multi-agent + twin + optimization + governance gate	The complete loop; D vs C isolates governance + optimization effect

Table 36 - Four-baseline experimental design.

24.3 Metrics, ablations and paper plan

Metrics: recovery cost (vs feasible floor), time-to-recover (signal-in to plan-executed), service level over the 30-day horizon, policy violations (must be zero for D), decision-trace completeness, and human-attention cost (approval clicks per event).

Ablations: remove twin (D vs C); remove gate (D vs an ungoverned D); swap LLM tier (small vs large model) to test where reasoning quality matters; perturb scenario parameters to test overfitting.

Paper skeleton: Abstract, Introduction, Problem, Related Work, Research Gap, Architecture, Methodology, Experimental Setup, Results, Ablation, Safety/Governance, Limitations, Future Work, Conclusion.

Candidate titles: (1) "Governed Recovery: A Multi-Agent LLM System for Self-Healing Pharmaceutical Supply Chains"; (2) "RecoveryBench: A Replayable Benchmark for Supply-Chain Recovery Decisions by Agentic AI"; (3) "LLM Agents that Call Solvers Safely: Guardrails for Autonomous Supply-Chain Re-planning"; (4) "From Weak Signals to Executed Recovery: Closing the Agentic Supply-Chain Loop"; (5) "Resilience as an Objective: Optimizing Time-to-Recover under Human Oversight".

Venues (student-realistic first): institute symposium and national student conferences; IEEE ICRA/IROS workshops on agentic systems; AAAI/IJCAI workshop tracks; INFORMS Applied Analytics; journals: Computers & Industrial Engineering, Supply Chain Management: An International Journal.

25. Patent and IP Roadmap

This chapter turns Chapter 6's white-space hypotheses into an actionable, student-executable IP plan. Two disclaimers bound everything: none of this is legal advice, and nothing here asserts patentability - every direction requires a professional prior-art and patentability search before any filing decision.

#	Invention hypothesis	Technical mechanism (what would be claimed, at high level)	Prior-art risk to check
I1	Guardrailed solver-calling agent architecture	A policy-validation layer that intercepts every agent tool-call toward an optimization engine, enforcing parameter bounds, budget caps and risk-tier permissions before solver execution, with failure-degradation to advisory mode	C3 AI orchestration patent; generic guardrail literature (Pre-Exec Bench); agent-gateway products

#	Invention hypothesis	Technical mechanism (what would be claimed, at high level)	Prior-art risk to check
I2	Weak-signal-to-parameterized-scenario pipeline	A pipeline that fuses multi-source weak signals into confidence-scored, schema-clamped disruption scenario parameters executable by a deterministic simulator, with analog-case grounding	Wagenitz risk-to-DES bridge; scenario-generation patents (national-level)
I3	GDP-compliance-aware cold-chain response	A compliance state machine coupling temperature telemetry to auto-generated investigation records, quarantine logic, and re-conditioning/re-booking actions	Tive/Controlant/Sensitech telemetry patents; excursion-management workflows
I4	Resilience-score-as-objective re-planning	A multi-objective re-planning formulation maximizing a defined composite resilience score (TTR, fill rate, cost-at-risk) subject to cost/policy constraints, with score provenance	Ivanov TTR/TTS methodology; planning-solver patents (Blue Yonder)
I5	Decision-audit ledger for agent systems	An append-only ledger capturing event-reasoning-scenario-approval-action-outcome chains with replay and threshold-recalibration hooks	Blockchain audit-log patents; SIEM/audit products

Table 37 - Five invention directions requiring professional review.

25.1 Student-executable filing strategy

Step 1 (now): treat the architecture as confidential until a provisional is filed - demo under NDA where possible, and keep claim-relevant implementation details out of public repos until Step 3.

Step 2 (post-finale, within weeks): engage a registered patent agent (India: provisional filing is modest-cost; many college IPR cells, incubators and SAP/college innovation grants subsidize). File a provisional on I1 + I2 combined (the strongest, most demo-aligned pairing); I3 remains the second-year candidate once pilot data exists.

Step 3 (12-month window): use the provisional year for pilot evidence and the professional prior-art search; convert to a full application only if the search and the business case both hold. Abandon gracefully if the white space closes - IP is a portfolio decision, not a trophy.

Parallel assets: even without filing, the ledger data, benchmark and calibration methods are trade-secret-able; the repo shows integrations, not the moat internals.

26. Startup Plan and Business Model

26.1 Company concept and ICP

Company concept: the post-alert decision layer for pharma supply chains - sold as "the system that turns disruption alerts into approved, executed, audited recovery." The ideal customer profile is an Indian pharmaceutical or life-sciences exporter with \$50m-500m revenue, an SAP or mid-market ERP backbone, 3-10 planning staff, regulated-market exposure (EU/US GDP), and a live scar from a disruption (a ruined cold-chain shipment, a demurrage bill, a missed tender). The buyer is the VP Supply Chain or COO; the user is the planning team lead; the compliance co-signer is QA/GDP. This segment is reachable (industry associations, pharma clusters in Hyderabad/Ahmedabad, SAP's own mid-market ecosystem), painful enough to buy, and small enough that giants will not custom-serve it first - the classic beachhead logic.

26.2 Business Model Canvas

Customer segments	Primary: Indian pharma exporters (\$50-500m). Secondary (Y2+): specialty chemicals, medical devices, auto components - all disruption-exposed, SAP-dense.
Value propositions	Minutes-not-weeks governed recovery; one prevented incident pays the year; GDP-grade audit trail for free; resilience score as a board-reportable KPI.
Channels	SAP ecosystem (hackfest credibility, SAP Store listing path, partner-manager intros), pharma industry associations, college incubator networks, outbound to SC leaders on LinkedIn.
Customer relationships	High-touch pilot to start (co-deploy with planning team), then quarterly business reviews with resilience scorecards; community of Indian SC planners for referrals.
Revenue streams	Tiered SaaS \$18k / \$36k / \$60k ACV (by sites + lanes + event volume) + success fee (5-10% of verified prevented-loss value, capped) + onboarding fee. Expansion: modules (ColdSentry, AlertTriage) and API access.
Key resources	The decision-outcome ledger (proprietary calibration data), RecoveryBench harness, solver-guardrail IP, SAP-integration expertise, the team's research credibility.
Key activities	Agent/twin/solver engineering, signal-source curation, compliance rulebook maintenance (per market), pilot delivery, benchmark publication.
Key partners	SAP (startup programs, BTP), cold-chain IoT vendors (telemetry-in), freight forwarders (execution-in), patent counsel, college IPR cell, pharma industry body for credibility.
Cost structure	Team salaries (largest), cloud + LLM inference (modest, usage-based), signal APIs, legal/IP, sales travel to pharma clusters. Gross margin target 75-85%.

Table 38 - Business Model Canvas (enterprise SaaS, mid-market first).

26.3 Go-to-market and the 1/3/5-year roadmap

GTM motion: paid pilot (\$8k-15k, 8-10 weeks, one network, success criteria pre-written) - convert to annual SaaS. Pilot sourcing: hackfest network, SAP partner intros, two warm pharma-cluster anchors.

The first-pilot playbook is written into Chapter 21's timeline.

Year 1: 2 paid pilots, 1 conversion; provisional filed; RecoveryBench released; 1 paper submitted.

Revenue: modest pilot income; the goal is proof, not profit.

Year 3: 10-14 customers in the beachhead; \$1.5-3.5m ARR band (base case); module expansion; first EU-market compliance rulebook; team of 8-12; seed round optional (founders keep control if base case holds).

Year 5: category position as "the recovery layer" across life-sciences verticals; \$6-12m ARR band (base); regional expansion (SEA generics exporters face identical dynamics); the ledger becomes the dataset competitors cannot synthesize.

26.4 "If the Hackfest ended tomorrow, what could we sell next week?"

The smallest commercially viable product is already visible in Layer 1: a Resilience Audit-as-a-Service. We take one exporter's network snapshot, run the twelve-scenario benchmark against it, and deliver a board-ready report: stockout probabilities per SKU-market under six disruption families, cost-at-risk quantified, and a prioritized resilience investment list. It needs no autonomous execution, no new integrations beyond a data extract, and it prices at \$3k-6k - a consulting-style wedge that funds the SaaS journey, seeds the ledger with real networks, and creates the pilot pipeline. The autonomous loop sells after trust is established; the audit sells on day one.

26.5 Moat ranking

Moat	Strength	Time to build	Defensibility note
Decision-outcome ledger (what worked, at what cost)	H	18-24 months	Generated only by running the loop on real networks; competitors cannot synthesize it
Solver-guardrail IP + patent portfolio	M-H	12-18 months	Deters feature-copying by alert vendors; requires professional review (Ch 18)
Compliance rulebooks per market (GDP/tariff)	M	12 months	Unsexy, cumulative, locally validated - enterprise buyers pay for exactly this

Moat	Strength	Time to build	Defensibility note
Benchmark + research credibility	M	6-12 months	RecoveryBench positions the team as the category's measurement standard
SAP integration depth	M	Ongoing	API-faithful connectors and A2A/MCP alignment; slows generic rivals more than SAP itself
Switching costs (ledger + rulebooks live in-tenant)	M	Per customer	Data gravity grows with every event processed

Table 39 - Moat analysis (strength: H/M/L).

27. Financial Model: Conservative, Base, Aggressive

Everything below is a model, not a forecast: the numbers instantiate the plan of Chapter 19 under three assumption sets, using benchmark-anchored inputs (SaaS median ACV \$24,266 per SaaS Capital 2026; mid-market bands \$15k-50k; gross margins 75-85% typical for mid-market SaaS with modest inference costs). Revenue figures are annual recurring revenue at year-end. The honest reading: Year-1 outcomes are dominated by execution, not by the model - the model's job is to prove the ceiling is worth the climb.

Assumption	Conservative	Base	Aggressive
ACV (blended, \$k)	18	30	45
New customers / yr (Y1; Y2; Y3; Y4; Y5)	1; 2; 4; 7; 10	2; 4; 8; 12; 16	3; 7; 14; 22; 30
Gross margin	78%	82%	85%
Fixed opex (\$k/yr, Y1 to Y5)	40; 90; 180; 320; 480	60; 150; 320; 560; 850	90; 240; 520; 900; 1400
Success-fee share of revenue	0%	8%	15%

Table 40 - Model assumptions (labeled ASSUMPTIONS; sources for anchors in text).




Figure 10 - ARR trajectories by scenario (USD millions, year-end). Base case crosses operating profitability in Year 4; aggressive in Year 3.

Metric (year-end)	Y1-C / B / A	Y3-C / B / A	Y5-C / B / A
Customers	1 / 2 / 3	7 / 14 / 24	24 / 42 / 76
ARR (\$m)	0.03 / 0.05 / 0.09	0.34 / 0.85 / 1.80	1.10 / 3.60 / 9.20
Gross profit (\$m)	0.02 / 0.04 / 0.08	0.27 / 0.70 / 1.53	0.86 / 2.95 / 7.82
Operating result (\$m)	-0.38 / -0.56 / -0.82	-0.09 / +0.38 / +1.01	+0.38 / +2.10 / +6.42

Table 41 - Model outputs (\$ = USD). C/B/A = Conservative/Base/Aggressive.

27.1 Customer ROI equation (the value model that prices the product)

The customer's annual value pool, with the equations we defend in Q&A: $V = V_{\text{stockout}} + V_{\text{expedite}} + V_{\text{excursion}} + V_{\text{inventory}} + V_{\text{labor}}$. $V_{\text{stockout}} = (P_{\text{incidents}} \times L_{\text{days}} \times \text{daily_margin_at_risk} \times \text{reduction_rate})$; $V_{\text{expedite}} = (\text{expedited_spend} \times \text{reduction_rate, benchmark } 15\text{-}25\%)$; $V_{\text{excursion}} = (\text{excursion_writeoffs} \times \text{prevention_rate, benchmark } 30\text{-}60\% \text{ with governed response})$; $V_{\text{inventory}} = (\text{buffer_inventory_value} \times \text{carrying_cost_rate} \times \text{optimization_gain, benchmark } 5\text{-}10\%)$; $V_{\text{labor}} = (\text{planner_hours_per_event} \times \text{events} \times \text{hourly_cost} \times \text{automation_rate, benchmark } 70\text{-}90\%)$. Worked example for a \$180m exporter (illustrative ASSUMPTIONS): 4 major incidents/yr x 6 days x \$220k daily margin-at-risk x 35% reduction = \$1.85m; expedite reduction on \$2.4m spend x 20% = \$0.48m; excursion write-offs \$0.6m x 45% = \$0.27m; buffer-carrying savings \$0.19m; labor savings \$0.09m - a value pool near \$2.9m against a \$30k-60k price. Even at one-tenth attribution discipline, the ROI case clears. Figure 11 shows the pool composition we present on stage.




Figure 11 - Composition of the annual customer value pool (worked example, shares of total). Every share is an ASSUMPTION to be replaced by pilot actuals.

28. Evidence Library and Glossary

Every statistic and claim in this unified dossier traces to public sources. The source index below consolidates the principal references from both source volumes; the selected references that follow list the works cited in the research deep-dive in APA style; and the glossary at the end defines the terms and acronyms used throughout. Claims sourced from social-media scans are directional color, never weighted evidence; numbers that could not be verified across at least one credible source were flagged in the source text rather than fabricated.

The dossier draws on 91 structured web searches whose raw results are archived as JSON in the project research directory. The index below lists the principal sources behind the report's major claims, with publication dates as carried in the search metadata. URLs are abbreviated to registrable domains where the full path appears in the chapter text.

#	Source and date	Used for
1	Gartner press release, Apr 7, 2026	Agentic SCM spend under \$2bn (2025) to \$53bn (2030); adoption 5% to 60%
2	Gartner via digitalapplied.com, Jun 2025 reading	Over 40% of agentic projects cancelled by end-2027
3	J.S. Held study via netsuite.com, Apr 16, 2026	\$184bn annual disruption cost
4	SAP news.sap.com, May 14 + May 22, 2026	Autonomous Supply Chain Management; sense-decide-act framing

#	Source and date	Used for
5	SAP news.sap.com, Oct 9, 2025; savictech.com Apr-May 2026	Joule agents wave (14 agents, GA Q1 2026); 6 assistants + 60+ agents at Sapphire 2026
6	help.sap.com (multiple pages, 2024-2026)	/IBP/API_SCENARIO; API_FREIGHTUNIT OData V4; Ariba Risk Exposure API (ARI-19205); HANA KG engine GA Q1 2025; GenAI Hub model catalog; BTP trial vs free tier
7	architecture.learning.sap.com, Apr 23, 2026	Build AI Agents on BTP; A2A + MCP protocols
8	techzine.eu, Nov 4, 2025	SAP TechEd 2025 MCP support
9	arXiv:2503.10009 (Mar 2025, v3 Aug 2025)	OR-LLM-Agent: 100% pass / 85% accuracy LLM-to-solver agent
10	arXiv, Jan 14, 2026	First synthesized SC disruption-monitoring evaluation dataset (30 scenarios)
11	Jannelli et al., Taylor & Francis, 2026	Agentic LLMs in the supply chain (autonomous consensus)
12	Tsinghua FIB lab, Nature-family survey, 2024	LLM-empowered agent-based modeling and simulation
13	Ivanov and Dolgui (2019-2025 canon); MDPI 2025 review	Digital supply chain twins; ripple effect; TTR/TTS
14	Wagenitz et al., SCS-Europe	LLM risk detection bridged to discrete-event simulation
15	Everstream (\$50m, Apr 2023); Prewave (\$67m, Jun 2024); Interos (~\$310m; \$40m Oct 2024 + \$20m Jan 2026); Altana (\$200m @ \$1bn, Jul 2024)	Competitor funding facts
16	TechMonitor / BlackKite / Interos, Nov-Dec 2024	Blue Yonder ransomware; 3.5m downstream estimate
17	Xeneta, Oct 4, 2024; project44, Sep 23, 2025	Port strike backlog weeks; Panama recovery mid-2025
18	ITF-OECD, Mar 18, 2024; FAO/OpenKnowledge (Drewry), Jan 2024; UNCTAD 2024	Cape re-routing costs; +417% freight spike; Red Sea transit collapse
19	whitehouse.gov, Jul 30, 2025; z2data, Dec 15, 2025; McKinsey, Dec 2025; NFTC, Aug 2025; KPMG, May 2025	De minimis end (Aug 29, 2025); 16.8% avg tariff; 30% of SC activities affected; 50% rerouting; 77% adapting
20	AlixPartners via CNBC, May 2021; JPMorgan, Apr 2023	Chip shortage \$110-210bn; auto production -26%
21	Chatterjee, PMC 2020; drugpatentwatch, Jan 20, 2026; CFR, Jun 3, 2026; ORF, Dec 10, 2025	India API/KSM dependency 70-90%; pharma choke point
22	Actizain, Mar 14, 2026	India pharma exports \$27.9bn FY2023-24
23	DPIIT study / PIB, Nov 27, 2025	India logistics cost 7.97% of GDP; NLP 5% target
24	Sensos, Jun 2025; tempcontrolpack, Nov 2025; pharmaceuticalmanufacturer.media, Feb 2026	\$35bn cold-chain losses; 20% biologics losses; 50% vaccines compromised
25	EMA.europa.eu; plumsense.com, Jun 2026	GDP obligations; excursion investigation requirements

#	Source and date	Used for
26	IBM think, May 29, 2025; Splunk, Feb 2025; StrangeBee (Gartner), Sep 2025	Alert fatigue and false positives
27	SaaS Capital, Aug 14, 2026; OpenView; RevTek	ACV benchmarks: median \$24,266; \$15k-50k mid-market band
28	Mordor Intelligence, Jan 23, 2026; Fortune Business Insights, Aug 2026; Kearney	SCRM market \$5-8bn, 9-14% CAGR
29	Dataminr (vendor data)	80% of manufacturers disrupted in 12-18 months; profit hits up to 13%
30	events.sap.com (South & East 2026); college360.co.in; valantic.com, 2025	Hackfest culture; business-value-first judging signals
31	patents.google.com: US12061673B1; WO2020142548A1; US20210374632A1; US 12,456,092; C3 AI IR, Oct 31, 2024	Prior-art screening set
32	github.com: amzn/supply-chain-simulation-environment; usnistgov/DiscreteEventLogisticsSystems; langgraph; pyomo; highs; or-tools	Open-source build inventory
33	research.google blog, Aug 31, 2026; Hugging Face model cards	TimesFM-3 zero-shot SOTA; Chronos/Moirai availability
34	newyorkfed.org	GSCPI macro pressure index
35	Kaggle datasets	Supply Chain Disruption and Recovery Dataset

Table A.1 - Principal sources.

Selected references

- Awasthi, R., et al. (2022). VacSIM: Learning effective strategies for COVID-19 vaccine distribution. ScienceDirect.
- Bain & Company. (2025, February 8). A roadmap for making India a global pharma exports hub.
- CDSCO. (2024, April 2). Guidelines on Good Distribution Practices for Pharmaceutical Products (CDSCO/GDP.PP Ver. 00).
- Cisco / Splunk. (2026, June 23). Cisco Time Series Model 1.0 for Gen-AI anomaly detection. Hugging Face.
- EROAD. (2025). The cost of temperature excursions in the pharmaceutical supply chain.
- European Commission. (2013). Guidelines of 5 November 2013 on Good Distribution Practice of medicinal products for human use (2013/C 343/01).
- Grand View Research. (2025). Cold chain monitoring market size, share and trends, 2024-2030.
- IBM Institute for Business Value. (2025, April 8). Scaling supply chain resilience: Agentic AI for autonomous operations.
- Jannelli, V. (2026). Agentic LLMs in the supply chain: Towards autonomous consensus-seeking.
- Kanwal, S. (2026). Temperature excursion risk window in pharmaceutical air cargo. Research Square.

- Konovalenko, I., & Ludwig, A. (2021). Comparison of machine learning classifiers for temperature monitoring in a pharmaceutical supply chain. ScienceDirect.
- Kumar, N. (2016). Temperature excursion management: A novel approach in pharmaceutical quality system. PMC.
- Mordor Intelligence. (2025, July 22). Healthcare cold chain logistics market size, share and 2030 growth.
- Mutambik, I. (2026). A blockchain-IoT-ML framework for sustainable vaccine cold chains. MDPI.
- Press Information Bureau, Government of India. (2025, December 17). India's pharmaceutical exports at USD 30.47 billion in 2024-25.
- SAP News. (2025, October 7). Innovative updates for supply chain management: Three new Joule agents.
- SAP News. (2026, May 14). Moving toward a more autonomous supply chain. SAP Sapphire 2026.
- SAP News. (2026, July 23). Self-healing supply chains may be closer than you think.
- Salesforce AI Research. (2024, November 11). GIFT-Eval: A benchmark for general time series forecasting. arXiv.
- UNDP India. (2021). Return on investment of the Electronic Vaccine Intelligence Network (eVIN). PMC.
- USP. (2022). General Chapter <1079.4> Temperature mapping for the qualification of storage areas.
- WHO. (2019). TRS 961, Annex 9: Model guidance for the storage and transport of time- and temperature-sensitive pharmaceutical products.

Term	Meaning in this dossier
Agent (AI)	A software process that perceives state, plans, and takes actions via tools under a policy; here, a specialized node in a LangGraph state machine
A2A / MCP	Agent2Agent protocol and Model Context Protocol - interoperability standards for agent communication and tool/data access; SAP-endorsed (2025-2026)
ACV	Annual contract value - the yearly revenue of a subscription contract
BTP	SAP Business Technology Platform - SAP's extension and integration platform (free tier used here)
CAP	SAP Cloud Application Programming Model - the framework used to build the API-faithful mock services
DE Mosul / DES	Discrete-event simulation - deterministic replay of network flows over time (SimPy)
Digital twin	A computable replica of the supply network used to replay scenarios before committing decisions
FEU	Forty-foot equivalent unit - container size used in freight-rate quotes
GDP	Good Distribution Practice - pharmaceutical distribution quality rules (EU/WHO/US) requiring documented excursion investigation
GraphRAG	Retrieval-augmented generation grounded in a knowledge graph plus vector search (SAP HANA pattern)
GSCPI	Global Supply Chain Pressure Index (New York Fed) - monthly macro disruption pressure measure

Term	Meaning in this dossier
HITL	Human-in-the-loop - risk-tiered human approval gates inside the autonomous loop
HS code	Harmonized System commodity code - the key on which tariff schedules apply
IBP	SAP Integrated Business Planning - demand/S&OP;/response/inventory planning suite

Term	Meaning in this dossier
KSM / API	Key starting material / active pharmaceutical ingredient - the upstream chemical inputs (India: 70-90% China dependency)
LangGraph	MIT-licensed agent orchestration framework with checkpointing and native human-in-the-loop interrupts
MILP	Mixed-integer linear programming - the optimization class used for re-planning (Pyomo + HiGHS)
OData	Open Data Protocol - REST-based API standard used by SAP products (IBP/TM shapes mocked with identical contracts)
Resilience score	Composite measure (time-to-recover, fill rate, cost-at-risk) used as an audited objective in re-planning
RTTVP	Real-time transportation visibility platform - the project44/FourKites category
SCRM	Supply chain risk management (software market category)
TTR / TTS	Time-to-recover / time-to-survive - Ivanov-school resilience metrics
Zero-shot forecasting	Time-series prediction by pretrained foundation models without per-customer training (Chronos, TimesFM)

Table B.1 - Glossary of key terms and acronyms.